
PyEpics: Python Epics Channel Access

Release 3.2.0

Matthew Newville

January 19, 2012

CONTENTS

1	Downloading and Installation	3
1.1	Prerequisites	3
1.2	Downloads	3
1.3	Testing	3
1.4	Development Version	4
1.5	Installation	4
1.6	Getting Started, Setting up the Epics Environment	4
1.7	Acknowledgments	5
1.8	Epics Open License	5
2	Overview of EPICS Channel Access in Python	7
2.1	Quick Start	8
2.2	Functions defined in <code>epics</code> : <code>caget()</code> , <code>caput()</code> , etc.	9
2.3	Motivation: Why another Python-Epics Interface?	12
2.4	Status and To-Do List	13
3	PV: Epics Process Variables	15
3.1	The PV class	15
3.2	String representation for a PV	20
3.3	Automatic Monitoring of a PV	20
3.4	User-supplied Callback functions	21
3.5	User-supplied Connection Callback functions	22
3.6	Put with wait, put callbacks, and <code>put_complete</code>	23
3.7	Examples	24
4	ca: Low-level Channel Access module	29
4.1	General description, difference with C library	29
4.2	Initialization, Finalization, and Life-cycle	30
4.3	Using the CA module	30
4.4	Implementation details	38
4.5	User-supplied Callback functions	39
4.6	Omissions	40
4.7	<code>CAThread</code> class	41
4.8	Examples	41
5	Devices: collections of PVs	45
5.1	Overview	45
5.2	Epics Motor Device	46
5.3	Other Device Examples	50

6	Alarms: respond when a PV goes out of range	53
6.1	Overview	53
6.2	Alarm Example	54
7	Auto-saving: simple save/restore of PVs	55
7.1	Overview	55
7.2	Examples	55
8	wx: wxPython Widgets for Epics	57
8.1	PV-aware Widgets	57
8.2	Decorators and other Utility Functions	60
8.3	wxMotorPanel Widget	61
8.4	OGL Classes	63
9	Advanced Topic with Python Channel Access	65
9.1	Working with waveform / array data	65
9.2	Strategies for working with large arrays	66
9.3	Using Python Threads	68
9.4	The wait and timeout options for get(), ca.get_complete()	71
9.5	Strategies for connecting to a large number of PVs	72
9.6	time.sleep() or epics.poll()?	73
	Python Module Index	75
	Index	77

PyEpics is an interface for the Channel Access (CA) library of the [Epics Control System](#) to the Python Programming language. The pyepics package provides a base `epics` module to python, with methods for reading from and writing to Epics Process Variables (PVs) via the CA protocol. The package includes a fairly complete, thin layer over the low-level Channel Access library in the `ca` module, and higher-level abstractions built on top of this basic functionality.

The package includes a simple, functional approach to CA similar to EZCA and the Unix command-line tools with functions in the main `epics` package including `epics.caget()`, `epics.caput()`, `epics.cainfo()`, and `epics.camonitor()`. There is also a `pv.PV` object which represents an Epics Process Variable as an easy-to-use Python object. Additional modules provide even higher-level programming support to Epics. These include groups of related PVs in `device.Device`, a simple method to create alarms in `alarm.Alarm`, and support for saving PVs values in the `autosave` module. Finally, there is support for conveniently tying epics PVs to wxPython widgets in the `wx` module.

In addition to the Pyepics library described here, several applications built with pyepics are available at <http://github.com/pyepics/epicsapps/>. See <http://pyepics.github.com/epicsapps/> for further details.

DOWNLOADING AND INSTALLATION

1.1 Prerequisites

This package requires Python version 2.5, 2.6, 2.7, or 3.2. It is possible that Python 2.4 will work if the ctypes package is installed, but this has not been tested. It should work with Python 3.1, as well, but this is no longer being tested as of this writing.

In addition, version 3.14 of the EPICS Channel Access library (v 3.14.8 or higher, I believe) is required. More specifically, the shared libraries libCom.so and libca.so (or Com.dll and ca.dll on Windows) from *Epics Base* are required to use this module. Using version 3.14.12 or higher is recommended – some of the features for ‘subarray records’ will only work with this 3.14.12 and higher.

For 32-bit Windows, pre-built DLLs from 3.14.12 (patched as of March, 2011) are included and installed so that no other Epics installation is required to use the modules. For Unix-like systems, these are assumed to be available (and findable by Python at runtime) on the system. This may mean setting LD_LIBRARY_PATH or DYLD_LIBRARY_PATH or configuring ldconfig.

The Python `numpy` module is not required, but will be used to convert EPICS waveforms values into numerical array data if available, and its use is strongly encouraged.

1.2 Downloads

The latest stable version of the PyEpics Package is 3.1.4. There are a few ways to get the PyEpics Package:

Download Option	Location
Source Kit	pyepics-3.1.4.tar.gz (CARS) or pyepics-3.1.4.tar.gz (PyPI)
Windows Installers	pyepics-3.1.4.win32-py2.6.exe or pyepics-3.1.4.win32-py2.7.exe or pyepics-3.1.4.win32-py3.2.exe
Development Version	use pyepics github repository

If you have `Python Setup Tools` installed, you can download and install the PyEpics Package simply with:

```
easy_install -U pyepics
```

1.3 Testing

Some automated unit-testing is done, using the tests folder from the source distribution kit. The following systems were tested for 3.1.4, all with Epics base 3.14.12.1. Except as noted, all tests pass. Those tests that fail are generally

well-understood.

Host OS	Epics HOST ARCH	Python	Failures, Notes
Linux	linux-x86	2.5.1	all pass
Linux	linux-x86	2.6	all pass
Linux	linux-x86	2.6.6	all pass
Linux	linux-x86_64	2.7	all pass
Linux	linux-x86_64	3.2	autosave fails
Mac OSX	darwin-x86	2.6.5	all pass
Windows	win32-x86	2.6.5	all pass
Windows	win32-x86	2.7.1	all pass
Windows	win32-x86	3.2.2	autosave fails

Testing Notes:

1. tests involving subarrays are known to fail with Epics base earlier than 3.14.11.
2. The autosave module relies on the 3rd part extension pyparsing, which seems to not work correctly for Python3.
3. The wx module is not automatically tested.
4. CA is not yet working with 64-bit Python on 64-bit Windows. It *does* work with 32-bit Python on 64-bit Linux, and Channel Access is known to work with 64-bit Windows. The current status is: I can get the 64-bit ca.dll to load with 64-bit Python, but there seems to be some disagreement about the lengths of basic C data types (for example, does a double take 8 or 16 bytes). This is being investigated....

1.4 Development Version

The PyEpics module is still under active development, with enhancements and bug-fixes are being added frequently. All development is done through the [pyepics github repository](#). To get a read-only copy of the latest version, use one of:

```
git clone http://github.com/pyepics/pyepics.git
git clone git@github.com:pyepics/pyepics.git
```

Current and older source source kits, and Windows Installers can also be found at the [PyEpics Source Tree](#).

1.5 Installation

Installation from source on any platform is:

```
python setup.py install
```

For more details, especially about how to set paths for LD_LIBRARY_PATH or DYLD_LIBRARY_PATH on Unix-like systems, see the INSTALL file.

Again, if you have [Python Setup Tools](#) installed, you can download and install the PyEpics Package with:

```
easy_install -U pyepics
```

1.6 Getting Started, Setting up the Epics Environment

In order for PyEpics to work at correctly, it must be able to find and load the Channel Access dynamic library (*libca.so*, *libca.dylib*, or *ca.dll* depending on the system). This dynamic library needs to be found at runtime.

There are a few ways to specify how to find this library:

1. set the environmental variable `PYEPICS_LIBCA` to the full path of the dynamic library, for example:

```
> export PYEPICS_LIBCA=/usr/local/epics/base-3.14.12.1/lib/linux-x86/libca.so
```
2. set the environmental variables `EPICS_BASE` and `EPICS_HOST_ARCH` to point to where the library was built. For example:

```
> export EPICS_BASE=/usr/local/epics/base-3.14.12.1  
> export EPICS_HOST_ARCH=linux-x86
```

will find the library at `/usr/local/epics/base-3.14.12.1/lib/linux-x86/libca.so`.

3. Put the dynamic library somewhere in the Python path. A convenient place might be the same `site-packages/pyepics` library folder as the python package is installed.

Note, that For Windows users, the DLLs (`ca.dll` and `Com.dll`) are included in the installation kit, and automatically installed to where they can be found at runtime (following rule 3 above).

With the Epics library loaded, it will need to be able to connect to Epics Process Variables. Generally, these variables are provided by Epics I/O controllers (IOCs) that are processes running on some device on the network. If you're connecting to PVs provided by IOCs on your local subnet, you should have no trouble. If trying to reach further network, you may need to set the environmental variable `EPICS_CA_ADDR_LIST` to specify which networks to search for PVs.

1.7 Acknowledgments

PyEpics was originally written and is maintained by Matt Newville <newville@cars.uchicago.edu>. Important contributions to the library have come from Angus Gratton, at the Australian National University. Several other people have provided valuable additions, suggestions, or bug reports, which has greatly improved the quality of the library: Michael Abbott, Marco Cammarata, Craig Haskins, Pete Jemian, Andrew Johnson, Janko Kolar, Irina Kosheleva, Tim Mooney, Eric Norum, Mark Rivers, Friedrich Schotte, Mark Vigder, Steve Wasserman, and Glen Wright.

1.8 Epics Open License

The PyEpics source code, this documentation, and all material associated with it are distributed under the Epics Open License:

The epics python module was originally written by

Matthew Newville <newville@cars.uchicago.edu> CARS, University of Chicago

There have been several contributions from many others, notably Angus Gratton <angus.gratton@anu.edu.au>. See the Acknowledgements section of the documentation for a list of more contributors.

Except where explicitly noted, all files in this distribution are licensed under the Epics Open License.:

Copyright 2010 Matthew Newville, The University of Chicago. All rights reserved.

The epics python module is distributed subject to the following license conditions: SOFTWARE LICENSE AGREEMENT Software: epics python module

1. The “Software”, below, refers to the epics python module (in either source code, or binary form and accompanying documentation). Each licensee is addressed as “you” or “Licensee.”
 2. The copyright holders shown above and their third-party licensors hereby grant Licensee a royalty-free nonexclusive license, subject to the limitations stated herein and U.S. Government license rights.
 3. You may modify and make a copy or copies of the Software for use within your organization, if you meet the following conditions:
 1. Copies in source code must include the copyright notice and this Software License Agreement.
 2. Copies in binary form must include the copyright notice and this Software License Agreement in the documentation and/or other materials provided with the copy.
 4. You may modify a copy or copies of the Software or any portion of it, thus forming a work based on the Software, and distribute copies of such work outside your organization, if you meet all of the following conditions:
 1. Copies in source code must include the copyright notice and this Software License Agreement;
 2. Copies in binary form must include the copyright notice and this Software License Agreement in the documentation and/or other materials provided with the copy;
 3. Modified copies and works based on the Software must carry prominent notices stating that you changed specified portions of the Software.
 5. Portions of the Software resulted from work developed under a U.S. Government contract and are subject to the following license: the Government is granted for itself and others acting on its behalf a paid-up, nonexclusive, irrevocable worldwide license in this computer software to reproduce, prepare derivative works, and perform publicly and display publicly.
 6. WARRANTY DISCLAIMER. THE SOFTWARE IS SUPPLIED “AS IS” WITHOUT WARRANTY OF ANY KIND. THE COPYRIGHT HOLDERS, THEIR THIRD PARTY LICENSORS, THE UNITED STATES, THE UNITED STATES DEPARTMENT OF ENERGY, AND THEIR EMPLOYEES: (1) DISCLAIM ANY WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, TITLE OR NON-INFRINGEMENT, (2) DO NOT ASSUME ANY LEGAL LIABILITY OR RESPONSIBILITY FOR THE ACCURACY, COMPLETENESS, OR USEFULNESS OF THE SOFTWARE, (3) DO NOT REPRESENT THAT USE OF THE SOFTWARE WOULD NOT INFRINGE PRIVATELY OWNED RIGHTS, (4) DO NOT WARRANT THAT THE SOFTWARE WILL FUNCTION UNINTERRUPTED, THAT IT IS ERROR-FREE OR THAT ANY ERRORS WILL BE CORRECTED.
 7. LIMITATION OF LIABILITY. IN NO EVENT WILL THE COPYRIGHT HOLDERS, THEIR THIRD PARTY LICENSORS, THE UNITED STATES, THE UNITED STATES DEPARTMENT OF ENERGY, OR THEIR EMPLOYEES: BE LIABLE FOR ANY INDIRECT, INCIDENTAL, CONSEQUENTIAL, SPECIAL OR PUNITIVE DAMAGES OF ANY KIND OR NATURE, INCLUDING BUT NOT LIMITED TO LOSS OF PROFITS OR LOSS OF DATA, FOR ANY REASON WHATSOEVER, WHETHER SUCH LIABILITY IS ASSERTED ON THE BASIS OF CONTRACT, TORT (INCLUDING NEGLIGENCE OR STRICT LIABILITY), OR OTHERWISE, EVEN IF ANY OF SAID PARTIES HAS BEEN WARNED OF THE POSSIBILITY OF SUCH LOSS OR DAMAGES.
-

OVERVIEW OF EPICS CHANNEL ACCESS IN PYTHON

The python `epics` package consists of several function, modules, and classes to interact with EPICS Channel Access. The simplest approach uses the functions `caget()`, `caput()`, `cainfo()`, `camonitor()`, and `camonitor_clear()` within the top-level `epics` module. These functions are similar to the standard command line utilities and to the EZCA library interface, and are described in more detail below.

To use the `epics` package, import it with:

```
import epics
```

The main components of this module include

- functions `caget()`, `caput()`, `camonitor()`, `camonitor_clear()`, and `cainfo()` as described below.
- a `ca` module, providing the low-level Epics Channel Access library as a set of functions.
- a `PV` object, giving a higher-level interface to Epics Channel Access.
- a `Device` object: a collection of related PVs
- a `Motor` object: a mapping of an Epics Motor
- an `Alarm` object, which can be used to set up notifications when a PV's values goes outside an acceptable bounds.
- an `epics.wx` module that provides wxPython classes designed for use with Epics PVs.

If you're looking to write quick scripts or a simple introduction to using Channel Access, the `caget()` and `caput()` functions are probably where you want to start.

If you're building larger scripts and programs, using `PV` objects provided by the `pv` module is recommended. The `PV` class provides a Process Variable object that has both methods (including `get()` and `put()`) to read and change the PV, and attributes that are kept automatically synchronized with the remote channel. For larger applications, you may find the `Device` class helpful.

The lowest-level CA functionality is exposed in the `ca` module, and companion `dbf` module. While not necessary recommended for most use cases, this module does provide a fairly complete wrapping of the basic EPICS CA library. For people who have used CA from C or other languages, this module should be familiar and seem quite usable, if a little more verbose and C-like than using `PV` objects.

In addition, the `epics` package contains more specialized modules for alarms, Epics motors, and several other *devices* (collections of `PVs`), and a set of wxPython widget classes for using EPICS `PVs` with wxPython.

The `epics` package is targeted for use on Unix-like systems (including Linux and Mac OS X) and Windows with Python versions 2.5, 2.6, 2.7, and 3.1, and 3.2.

2.1 Quick Start

Whether you're familiar with Epics Channel Access or not, start here. You'll then be able to use Python's introspection tools and built-in help system, and the rest of this document as a reference and for detailed discussions.

2.1.1 Functional Approach: `caget()`, `caput()`

To get values from PVs, you can use the `caget ()` function:

```
>>> from epics import caget, caput
>>> m1 = caget('XXX:m1.VAL')
>>> print m1
1.2001
```

To set PV values, you can use the `caput ()` function:

```
>>> caput('XXX:m1.VAL', 1.90)
>>> print caget('XXX:m1.VAL')
1.9000
```

For many cases, this approach is ideal because of its simplicity and clarity.

2.1.2 Object Oriented Approach: `PV`

If you want to repeatedly access the same PV, you may find it more convenient to "create a PV object" and use it in a more object-oriented manner.

```
>>> from epics import PV
>>> pv1 = PV('XXX:m1.VAL')
```

PV objects have several methods and attributes. The most important methods are `get ()` and `put ()` to receive and send the PV's value, and the `value` attribute which stores the current value. In analogy to the `caget ()` and `caput ()` examples above, the value of a PV can be fetched either with

```
>>> print pv1.get()
1.2001
```

or

```
>>> print pv1.value
1.2001
```

To set a PV's value, you can either use

```
>>> pv1.put(1.9)
```

or assign the `value` attribute

```
>>> pv1.value = 1.9
```

You can see a few of the most important properties of a PV by simply printing it:

```
>>> print pv1
<PV 'XXX:m1.VAL', count=1, type=double, access=read/write>
```

Even more complete information can be seen by printing the PVs `info` attribute:

```
>>> print pv1.info
== XXX:m1.VAL (native_double) ==
  value      = 1.9
  char_value = '1.90000'
  count      = 1
  nelm       = 1
  type       = double
  units      = mm
  precision  = 5
  host       = somehost.cars.aps.anl.gov:5064
  access     = read/write
  status     = 0
  severity   = 0
  timestamp  = 1265996457.212 (2010-Feb-12 11:40:57.212)
  upper_ctrl_limit = 12.5
  lower_ctrl_limit = -12.3
  upper_disp_limit = 12.5
  lower_disp_limit = -12.3
  upper_alarm_limit = 0.0
  lower_alarm_limit = 0.0
  upper_warning_limit = 0.0
  lower_warning_limit = 0.0
  PV is internally monitored, with 0 user-defined callbacks:
=====
```

PV objects have several additional methods, especially related to monitoring changes to the PVs and defining functions to be run when the value does change. There are also attributes associated with a PVs *Control Attributes*, like those shown above in the `info` attribute. Further details are at [PV: Epics Process Variables](#).

2.2 Functions defined in `epics`: `caget()`, `caput()`, etc.

The simplest interface to EPICS Channel Access provides functions `caget()`, `caput()`, as well as functions `camonitor()`, `camonitor_clear()`, and `cainfo()`. These functions are similar to the EPICS command line utilities and to the functions in the EZCA library. They all take the name of an Epics Process Variable (PV) as the first argument. As with the EZCA library, the python implementation keeps an internal cache of connected PV (in this case, using *PV* objects) so that repeated use of a PV name will not actually result in a new connection. Thus, though the functionality is limited, the performance of the functional approach can be quite good.

2.2.1 `caget()`

```
epics.caget(pvname[, as_string=False[, count=None[, as_numpy=True[, timeout=None[,
    use_monitor=False]]]])
    retrieves and returns the value of the named PV.
```

Parameters

- **pvname** – name of Epics Process Variable
- **as_string** (True/False) – whether to return string representation of the PV value.
- **count** (integer or None) – number of elements to return for array data.
- **as_numpy** (True/False) – whether to return the Numerical Python representation for array data.

- **timeout** (float or None) – maximum time to wait (in seconds) for value before returning None.
- **use_monitor** (True/False) – whether to rely on monitor callbacks or explicitly get value now.

The *count* and *as_numpy* options apply only to array or waveform data. The default behavior is to return the full data array and convert to a numpy array if available. The *count* option can be used to explicitly limit the number of array elements returned, and *as_numpy* can turn on or off conversion to a numpy array.

The *timeout* argument sets the maximum time to wait for a value to be fetched over the network. If the timeout is exceeded, `caget()` will return None. This might imply that the PV is not actually available, but it might also mean that the data is large or network slow enough that the data just hasn't been received yet, but may show up later.

The *use_monitor* argument sets whether to rely on the monitors from the underlying PV. The default is `False`, so that each `caget()` will explicitly ask the value to be sent instead of relying on the automatic monitoring normally used for persistent PVs. If this makes no sense, leaving the default value of `True` is fine. For more details on making `caget()` more efficient, see *Automatic Monitoring of a PV* and *The wait and timeout options for get(), ca.get_complete()*.

The *as_string* argument tells the function to return the **string representation** of the value. The details of the string representation depends on the variable type of the PV. For integer (short or long) and string PVs, the string representation is pretty easy: 0 will become '0', for example. For float and doubles, the internal precision of the PV is used to format the string value. For enum types, the name of the enum state is returned:

```
>>> from epics import caget, caput, cainfo
>>> print caget('XXX:m1.VAL')      # A double PV
0.10000000000000001

>>> print caget('XXX:m1.DESC')     # A string PV
'Motor 1'

>>> print caget('XXX:m1.FOFF')     # An Enum PV
1
```

Adding the *as_string=True* argument always results in string being returned, with the conversion method depending on the data type:

```
>>> print caget('XXX:m1.VAL', as_string=True)
'0.10000'

>>> print caget('XXX:m1.FOFF', as_string=True)
'Frozen'
```

For most array data from Epics waveform records, the regular value will be a numpy array (or a python list if numpy is not installed). The string representation will be something like '<array size=128, type=int>' depending on the size and type of the waveform. An array of doubles might be:

```
>>> print caget('XXX:scan1.P1PA')  # A Double Waveform
array([-0.08      , -0.078      , -0.076      , ...,
       1.99599814, 1.99799919,  2.        ])

>>> print caget('XXX:scan1.P1PA', as_string=True)
'<array size=2000, type=DOUBLE>'
```

As an important special case, CHAR waveforms will be turned to Python strings when *as_string* is `True`. This is useful to work around the low limit of the maximum length (40 characters!) of EPICS strings, and means that it is fairly common to use CHAR waveforms when long strings are desired:

```
>>> print caget('XXX:dir')          # A CHAR waveform
array([ 84,  58,  92, 120,  97, 115,  95, 117, 115,
       101, 114,  92,  77,  97, 114,  99, 104,  50,  48,
```

```

49, 48, 92, 70, 97, 115, 116, 77, 97, 112])

>>> print caget('XXX:dir', as_string=True)
'T:\\xas_user\\March2010\\Fastmap'

```

Of course, character waveforms are not always used for long strings, but can also hold byte array data, such as comes from some detectors and devices.

2.2.2 caput ()

```
epics.caput (pvname, value[, wait=False[, timeout=60]])
```

set the value of the named PV.

Parameters

- **pvname** – name of Epics Process Variable
- **value** – value to send.
- **wait** (True/False) – whether to wait until the processing has completed.
- **timeout** (*double*) – how long to wait (in seconds) for put to complete before giving up.

Return type integer

The optional *wait* argument tells the function to wait until the processing completes. This can be useful for PVs which take significant time to complete, either because it causes a physical device (motor, valve, etc) to move or because it triggers a complex calculation or data processing sequence. The *timeout* argument gives the maximum time to wait, in seconds. The function will return after this (approximate) time even if the `caput ()` has not completed.

This function returns 1 on success, and a negative number if the timeout has been exceeded.

```

>>> from epics import caget, caput, cainfo
>>> caput ('XXX:m1.VAL', 2.30)
1
>>> caput ('XXX:m1.VAL', -2.30, wait=True)
... waits a few seconds ...
1

```

2.2.3 cainfo()

```
epics.cainfo (pvname[, print_out=True])
```

prints (or returns as a string) an informational paragraph about the PV, including Control Settings.

Parameters

- **pvname** – name of Epics Process Variable
- **print_out** – whether to write results to standard output (otherwise the string is returned).

2.2.4 camonitor()

```
epics.camonitor (pvname[, writer=None[, callback=None]])
```

This *sets a monitor* on the named PV, which will cause *something* to be done each time the value changes. By default the PV name, time, and value will be printed out (to standard output) when the value changes, but the action that actually happens can be customized.

Parameters

- **pvname** – name of Epics Process Variable
- **writer** (`None` or a callable function that takes a string argument.) – where to write results to standard output .
- **callback** (`None` or callable function) – user-supplied function to receive result

One can specify any function that can take a string as *writer*, such as the `write()` method of an open file that has been open for writing. If left as `None`, messages of changes will be sent to `sys.stdout.write()`. For more complete control, one can specify a *callback* function to be called on each change event. This callback should take keyword arguments for *pvname*, *value*, and *char_value*. See [User-supplied Callback functions](#) for information on writing callback functions for `camonitor()`.

```
>>> from epics import camonitor
>>> camonitor('XXX.m1.VAL')
XXX.m1.VAL 2010-08-01 10:34:15.822452 1.3
XXX.m1.VAL 2010-08-01 10:34:16.823233 1.2
XXX.m1.VAL 2010-08-01 10:34:17.823233 1.1
XXX.m1.VAL 2010-08-01 10:34:18.823233 1.0
```

2.2.5 `camonitor_clear()`

`epics.camonitor_clear(pvname)`
clears a monitor set on the named PV by `camonitor()`.

Parameters *pvname* – name of Epics Process Variable

This simple example monitors a PV with `camonitor()` for while, with changes being saved to a log file. After a while, the monitor is cleared and the log file is inspected:

```
>>> import epics
>>> fh = open('PV1.log', 'w')
>>> epics.camonitor('XXX:DMM1Ch2_calc.VAL', writer=fh.write)
>>> .... wait for changes ...
>>> epics.camonitor_clear('XXX:DMM1Ch2_calc.VAL')
>>> fh.close()
>>> fh = open('PV1.log', 'r')
>>> for i in fh.readlines(): print i[:-1]
XXX:DMM1Ch2_calc.VAL 2010-03-24 11:56:40.536946 -183.5035
XXX:DMM1Ch2_calc.VAL 2010-03-24 11:56:41.536757 -183.6716
XXX:DMM1Ch2_calc.VAL 2010-03-24 11:56:42.535568 -183.5112
XXX:DMM1Ch2_calc.VAL 2010-03-24 11:56:43.535379 -183.5466
XXX:DMM1Ch2_calc.VAL 2010-03-24 11:56:44.535191 -183.4890
XXX:DMM1Ch2_calc.VAL 2010-03-24 11:56:45.535001 -183.5066
XXX:DMM1Ch2_calc.VAL 2010-03-24 11:56:46.535813 -183.5085
XXX:DMM1Ch2_calc.VAL 2010-03-24 11:56:47.536623 -183.5223
XXX:DMM1Ch2_calc.VAL 2010-03-24 11:56:48.536434 -183.6832
```

2.3 Motivation: Why another Python-Epics Interface?

PyEpics version 3 is intended as an improvement over EpicsCA 2.1, and should replace that older Epics-Python interface. That version had performance issues, especially when connecting to a large number of PVs, is not thread-aware, and has become difficult to maintain for Windows and Linux.

There are a few other Python modules exposing Epics Channel Access available. Most of these have a interface to the CA library that was both closer to the C library and lower-level than EpicsCA. Most of these interfaces use specialized C-Python ‘wrapper’ code to provide the interface.

Because of this, an additional motivation for this package was to allow a more common interface to be used that built higher-level objects (as EpicsCA had) on top of a complete lower-level interface. The desire to come to a more universally-acceptable Python-Epics interface has definitely influenced the goals for this module, which include:

1. providing both low-level (C-like) and higher-level access (Pythonic objects) to the EPICS Channel Access protocol.
2. supporting as many features of Epics 3.14 as possible, including preemptive callbacks and thread support.
3. easy support and distribution for Windows and Unix-like systems.
4. being ready for porting to Python3.
5. using Python’s ctypes library.

The main implementation feature used here (and difference from EpicsCA) is using Python’s ctypes library to handle the connection between Python and the CA C library. Using ctypes has many advantages. Principally, it fully eliminates the need to write (and maintain) wrapper code either with SWIG or directly with Python’s C API. Since the ctypes module allows access to C data and objects in pure Python, no compilation step is needed to build the module, making installation and support on multiple platforms much easier. Since ctypes loads a shared object library at runtime, the underlying Epics Channel Access library can be upgraded without having to re-build the Python wrapper. In addition, using ctypes provides the most reliable thread-safety available, as each call to the underlying C library is automatically made thread-aware without explicit code. Finally, by avoiding the C API altogether, migration to Python3 is greatly simplified. PyEpics3 does work with both Python 2.* and 3.*.

2.4 Status and To-Do List

The PyEpics package is under active development. The current status is that most features are working well, and it is starting to be used in production code, but more testing and better tests are needed.

The package is targeted and tested to work with Python 2.5, 2.6, 2.7, and 3.1 simultaneously (that is, the same code is meant to support all versions). Currently, the package works with Python 3.1, but is not extremely well-tested.

There are several desired features are left undone or unfinished:

- port CaChannel interface, ca_util, epicsPV (and other interfaces??) to use epics.ca
- add more “devices”, including low-level epics records.
- incorporate some or all of the Channel Access Server from [pcaspy](#)
- build and improve applications.

PV: EPICS PROCESS VARIABLES

The `pvs` module provides a higher-level class `pvs.PV`, which creates a *PV* object for an EPICS Process Variable. A *PV* object has both methods and attributes for accessing its properties.

3.1 The `PV` class

```
class pvs.PV(pvname[, callback=None[, form='native'[, auto_monitor=None[, connec-
            tion_callback=None[, connection_timeout=None[, verbose=False]]]]])
create a PV object for a named Epics Process Variable.
```

Parameters

- **pvname** – name of Epics Process Variable
- **callback** (*callable, tuple, list or None*) – user-defined callback function on changes to PV value or state.
- **form** (*string, one of ('native','ctrl', or 'time')*) – which epics *data type* to use: the 'native' , or the 'ctrl' (Control) or 'time' variant.
- **auto_monitor** (*None, True, False, or bitmask (see [Automatic Monitoring of a PV](#))*) – whether to automatically monitor the PV for changes.
- **connection_callback** (*callable or None*) – user-defined function called on changes to PV connection status.
- **connection_timeout** (*float or None*) – time (in seconds) to wait for connection before giving up
- **verbose** (*True/False*) – whether to print out debugging messages

Once created, a PV should (barring any network issues) automatically connect and be ready to use.

```
>>> from epics import PV
>>> p = PV('XX:m1.VAL')
>>> print p.get()
>>> print p.count, p.type
```

The *pvname* is required, and is the name of an existing Process Variable.

The *callback* parameter specifies one or more python methods to be called on changes, as discussed in more detail at [User-supplied Callback functions](#)

The *connection_callback* parameter specifies a python method to be called on changes to the connection status of the PV (that is, when it connects or disconnects). This is discussed in more detail at [User-supplied Connection Callback functions](#)

The *form* parameter specifies which of the three variants ‘native’ (the default), ‘ctrl’ (Control) or ‘time’ to use for the PV. The control and time variants add additional fields to the PV, which can be useful in some cases. Also note that the additional ‘ctrl’ value fields (see the [Table of Control Attributes](#)) can be obtained with `get_ctrlvars()` even for PVs of ‘native’ form.

The *auto_monitor* parameter specifies whether the PV should be automatically monitored. See [Automatic Monitoring of a PV](#) for a detailed description of this.

The *verbose* parameter specifies more verbose output on changes, and is intended for debugging purposes.

3.1.1 methods

A PV has several methods for getting and setting its value and defining callbacks to be executed when the PV changes.

`pv.get([count=None[, as_string=False[, as_numpy=True[, timeout=None[, use_monitor=True]]]])`
get and return the current value of the PV

Parameters

- **count** (integer or None) – maximum number of array elements to return
- **as_string** (True/False) – whether to return the string representation of the value.
- **as_numpy** – whether to try to return a numpy array where appropriate.
- **timeout** (float or None) – maximum time to wait for data before returning None.
- **use_monitor** (True/False) – whether to rely on monitor callbacks or explicitly get value now.

see [String representation for a PV](#) for details on how the string representation is determined.

With the *as_numpy* option, an array PV (that is, a PV whose value has more than one element) will be returned as a numpy array, provided the numpy module is available. See [Strategies for working with large arrays](#) for a discussion of strategies for how to best deal with very large arrays.

The *use_monitor* option controls whether the most recent value from the automatic monitoring will be used or whether the value will be explicitly asked for right now. Usually, you can rely on a PVs value being kept up to date, and so the default here is `True`. But, since network traffic is not instantaneous and hard to predict, the value returned with *use_monitor=True* may be out-of-date.

The *timeout* sets how long (in seconds) to wait for the value to be sent. This only applies with *use_monitor=False*, or if the PV is not automatically monitored. Otherwise, the most recently received value will be sent immediately.

See [Automatic Monitoring of a PV](#) for more on monitoring PVs and [The wait and timeout options for get\(\), ca.get_complete\(\)](#) for more details on what happens when a `pv.get()` times out.

`pv.put(value[, wait=False[, timeout=30.0[, use_complete=False[, callback=None[, callback_data=None]]]])`

set the PV value, optionally waiting to return until processing has completed, or setting the `put_complete` to indicate complete-ness.

Parameters

- **value** – value to set PV
- **wait** (True/False) – whether to wait for processing to complete (or time-out) before returning.
- **timeout** (float) – maximum time to wait for processing to complete before returning anyway.

- **use_complete** (True/False) – whether to use a built-in callback to set `put_complete`.
- **callback** (None or a valid python function) – user-supplied function to run when processing has completed.
- **callback_data** – extra data to pass on to a user-supplied callback function.

The *wait* and *callback* arguments, as well as the ‘use_complete’ / `put_complete` attribute give a few options for knowing that a `put ()` has completed. See *Put with wait, put callbacks, and put_complete* for more details.

`pv.get_ctrlvars ()`

returns a dictionary of the **control values** for the PV. This dictionary may have many members, depending on the data type of PV. See the *Table of Control Attributes* for details.

`pv.poll ([evt=1.e-4[, iot=1.0]])`

poll for changes. This simply calls `ca.poll ()`

Parameters

- **evt** (float) – time to pass to `ca.pend_event ()`
- **iot** (float) – time to pass to `ca.pend_io ()`

`pv.connect ([timeout=None])`

this explicitly connects a PV, and returns whether or not it has successfully connected. It is probably not that useful, as connection should happen automatically. See `wait_for_connection ()`.

Parameters **timeout** (float) – maximum connection time, passed to `ca.connect_channel ()`

Return type True/False

if timeout is None, the PVs `connection_timeout` parameter will be used. If that is also None, `ca.DEFAULT_CONNECTION_TIMEOUT` will be used.

`pv.wait_for_connection ([timeout=None])`

this waits until a PV is connected, or has timed-out waiting for a connection. Returns whether the connection has occurred.

Parameters **timeout** (float) – maximum connection time.

Return type True/False

if timeout is None, the PVs `connection_timeout` parameter will be used. If that is also None, `ca.DEFAULT_CONNECTION_TIMEOUT` will be used.

`pv.disconnect ()`

disconnect a PV, clearing all callbacks.

`pv.add_callback (callback=None[, index=None[, with_ctrlvars=True[, **kw]])`

adds a user-defined callback routine to be run on each change event for this PV. Returns the integer *index* for the callback.

Parameters

- **callback** (None or callable) – user-supplied function to run when PV changes.
- **index** (None (integer will be produced) or immutable) – identifying key for this callback
- **with_ctrlvars** – whether to (try to) make sure that accurate `control values` will be sent to the callback.
- **kw** – additional keyword/value arguments to pass to each execution of the callback.

Return type integer

Note that multiple callbacks can be defined, each having its own index (a dictionary key, typically an integer). When a PV changes, all the defined callbacks will be executed. They will be called in order (by sorting the keys of the `callbacks` dictionary)

See also: `callbacks` attribute, *User-supplied Callback functions*

`pv.remove_callback (index=None)`

remove a user-defined callback routine using supplied

Parameters `index` (None or integer) – index of user-supplied function, as returned by `add_callback()`, and also to key for this callback in the `callbacks` dictionary.

Return type integer

If only one callback is defined an `index=None`, this will clear the only defined callback.

See also: `callbacks` attribute, *User-supplied Callback functions*

`pv.clear_callbacks ()`

remove all user-defined callback routine.

`pv.run_callbacks ()`

execute all user-defined callbacks right now, even if the PV has not changed. Useful for debugging!

See also: `callbacks` attribute, *User-supplied Callback functions*

`pv.run_callback (index)`

execute a particular user-defined callback right now, even if the PV has not changed. Useful for debugging!

See also: `callbacks` attribute, *User-supplied Callback functions*

3.1.2 attributes

A PV object has many attributes, each associated with some property of the underlying PV: its *value*, *host*, *count*, and so on. For properties that can change, the PV attribute will hold the latest value for the corresponding property. Most attributes are **read-only**, and cannot be assigned to. The exception to this rule is the `value` attribute.

`pv.value`

The current value of the PV.

Note: The `value` attribute can be assigned to. When read, the latest value will be returned, even if that means a `get()` needs to be called.

Assigning to `value` is equivalent to setting the value with the `put()` method.

```
>>> from epics import PV
>>> p1 = PV('xxx.VAL')
>>> print p1.value
1.00
>>> p1.value = 2.00
```

`pv.char_value`

The string representation of the string, as described in `get()`.

`pv.status`

The PV status, which will be 1 for a Normal, connected PV.

`pv.type`

string describing data type of PV, such as *double*, *float*, *enum*, *string*, *int*, *long*, *char*, or one of the *ctrl* or *time* variants of these, which will be named *ctrl_double*, *time_enum*, and so on. See the *Table of DBR Types*

pv.ftype

The integer value (from the underlying C library) indicating the PV data type according to [Table of DBR Types](#)

pv.host

string of host machine provide this PV.

pv.count

number of data elements in a PV. 1 except for waveform PVs, where it gives the number of elements in the waveform. For recent versions of Epics Base (3.14.11 and later?), this gives the *.NORD* field, which gives the number of elements last put into the PV and which may be less than the maximum number allowed (see *nelm* below).

pv.nelm

number of data elements in a PV. 1 except for waveform PVs where it gives the maximum number of elements in the waveform. For recent versions of Epics Base (3.14.11 and later?), this gives the *.NELM* parameter. See also the *count* attribute above.

pv.read_access

Boolean (True/False) for whether PV is readable

pv.write_access

Boolean (True/False) for whether PV is writable

pv.access

string describing read/write access. One of 'read/write', 'read-only', 'write-only', 'no access'.

pv.severity

severity value of PV. Usually 0 for PVs that are not in an alarm condition.

pv.timestamp

Unix (not Epics!!) timestamp of the last seen event for this PV.

pv.precision

number of decimal places of precision to use for float and double PVs

pv.units

string of engineering units for PV

pv.enum_strs

a list of strings for the enumeration states of this PV (for enum PVs)

pv.info

a string paragraph (ie, including newlines) showing much of the information about the PV.

pv.upper_disp_limit**pv.lower_disp_limit****pv.upper_alarm_limit****pv.lower_alarm_limit****pv.lower_warning_limit****pv.upper_warning_limit****pv.upper_ctrl_limit****pv.lower_ctrl_limit**

These are all the various kinds of limits for a PV.

pv.put_complete

a Boolean (True/False) value for whether the most recent `put ()` has completed.

pv.callbacks

a dictionary of currently defined callbacks, to be run on changes to the PV. This dictionary has integer keys (generally in increasing order of when they were defined) which sets which order for executing the callbacks. The values of this dictionary are tuples of (*callback*, *keyword_arguments*).

Note: The **callbacks** attribute can be assigned to or manipulated directly. This is not recommended. Use the methods `add_callback()`, `remove_callback()`, and `clear_callbacks()` instead of altering this dictionary directly.

pv.connection_callbacks

a simple list of connection callbacks: functions to be run when the connection status of the PV changes. See *User-supplied Connection Callback functions* for more details.

3.2 String representation for a PV

The string representation for a *PV*, as returned either with the *as_string* argument to `ca.get()` or from the `char_value` attribute (they are equivalent) needs some further explanation.

The value of the string representation (hereafter, the `char_value`), will depend on the native type and count of a *PV*. *Table of String Representations*

Table of String Representations: How raw data `value` is mapped to `char_value` for different native data types.

<i>data types</i>	<i>count</i>	<i>char_value</i>
string	1	= value
char	1	= value
short	1	= str(value)
long	1	= str(value)
enum	1	= enum_str[value]
double	1	= ("%%.%if" % (precision)) % value
float	1	= ("%%.%if" % (precision)) % value
char	> 1	= long string from bytes in array
all others	> 1	= <array size=*count*, type=*type*>

For double/float values with large exponents, the formatting will be ("%%.%ig" % (*precision*)) % *value*. For character waveforms (*char* data with *count* > 1), the `char_value` will be set according to:

```
>>> firstnull = val.index(0)
>>> if firstnull == -1: firstnull= len(val)
>>> char_value = ''.join([chr(i) for i in val[:firstnull]].rstrip())
```

3.3 Automatic Monitoring of a PV

When creating a PV, the *auto_monitor* parameter specifies whether the PV should be automatically monitored or not. Automatic monitoring means that an internal callback will be registered for changes. Any callbacks defined by the user will be called by this internal callback when changes occur.

For most scalar-value PVs, this automatic monitoring is desirable, as the PV will see all changes (and run callbacks) without any additional interaction from the user. The PV's value will always be up-to-date and no unnecessary network traffic is needed.

Possible values for `auto_monitor` are:

False For some PVs, especially those that change much more rapidly than you care about or those that contain large arrays as values, `auto_monitoring` can add network traffic that you don't need. For these, you may wish to create your PVs with `auto_monitor=False`. When you do this, you will need to make calls to `get()` to explicitly get the latest value.

None The default value for `auto_monitor` is `None`, and is set to `True` if the element count for the PV is smaller than `ca.AUTOMONITOR_MAXLENGTH` (default of 65536). To suppress monitoring of PVs with fewer array values, you will have to explicitly turn `auto_monitor` to `False`. For waveform arrays with more elements, automatic monitoring will not be done unless you explicitly set `auto_monitor=True`, or to an explicit mask. See [Strategies for working with large arrays](#) for more details.

True When `auto_monitor` is set to `True`, the value will be monitored using the default subscription mask set at `ca.DEFAULT_SUBSCRIPTION_MASK`.

This mask determines which kinds of changes cause the PV to update. By default, the subscription updates when the PV value changes by more than the monitor deadband, or when the PV alarm status changes. This behavior is the same as the default in EPICS' `camonitor` tool.

Mask It is also possible to request an explicit type of CA subscription by setting `auto_monitor` to a numeric subscription mask made up of `dbt.DBE_ALARM`, `dbt.DBE_LOG` and/or `dbt.DBE_VALUE`. This mask will be passed directly to `ca.create_subscription()`. An example would be:

```
pv1 = PV('AAA', auto_monitor=dbt.DBE_VALUE)
pv2 = PV('BBB', auto_monitor=dbt.DBE_VALUE|dbt.DBE_ALARM)
pv3 = PV('CCC', auto_monitor=dbt.DBE_VALUE|dbt.DBE_ALARM|dbt.DBE_LOG)
```

which will generate callbacks for `pv1` only when the value of 'AAA' changes, while `pv2` will receive callbacks if the value or alarm state of 'BBB' changes, and `pv3` will receive callbacks for all changes to 'CCC'. Note that these `dbt.DBE_****` constants are Ored together as a bitmask.

3.4 User-supplied Callback functions

This section describes user-defined functions that are called when the value of a PV changes. These callback functions are useful as they allow you to be notified of changes without having to continually ask for a PV's current value. Much of this information is similar to that in [User-supplied Callback functions](#) for the `ca` module, though there are some important enhancements to callbacks on *PV* objects.

You can define more than one callback function per PV to be run on value changes. These functions can be specified when creating a PV, with the `callback` argument which can take either a single callback function or a list or tuple of callback functions. After a PV has been created, you can add callback functions with `add_callback()`, remove them with `remove_callback()`, and explicitly run them with `run_callback()`. Each callback has an internal unique *index* (a small integer number) that can be used for specifying which one to add, remove, and run.

When defining a callback function to be run on changes to a PV, it is important to know two things:

1. how your function will be called.
2. what is permissible to do inside your callback function.

Callback functions will be called with several keyword arguments. You should be prepared to have them passed to your function, and should always include `**kw` to catch all arguments. Your callback will be sent the following keyword parameters:

- *pvname*: the name of the pv
- *value*: the latest value
- *char_value*: string representation of value

- *count*: the number of data elements
- *ftype*: the numerical CA type indicating the data type
- *type*: the python type for the data
- *status*: the status of the PV (1 for OK)
- *precision*: number of decimal places of precision for floating point values
- *units*: string for PV units
- *severity*: PV severity
- *timestamp*: timestamp from CA server.
- *read_access*: read access (True/False)
- *write_access*: write access (True/False)
- *access*: string description of read- and write-access
- *host*: host machine and CA port serving PV
- *enum_strs*: the list of enumeration strings
- *upper_disp_limit*: upper display limit
- *lower_disp_limit*: lower display limit
- *upper_alarm_limit*: upper alarm limit
- *lower_alarm_limit*: lower alarm limit
- *upper_warning_limit*: upper warning limit
- *lower_warning_limit*: lower warning limit
- *upper_ctrl_limit*: upper control limit
- *lower_ctrl_limit*: lower control limit
- *chid*: integer channel ID
- ***cb_info*: (index, self) tuple containing callback ID and the PV object**

Some of these may not be directly applicable to all PV data types, and some values may be `None` if the control parameters have not yet been fetched with `get_ctrlvars()`.

It is important to keep in mind that the callback function will be run *inside* a CA function, and cannot reliably make any other CA calls. It is helpful to think “this all happens inside of a `pend_event()` call”, and in an epics thread that may or may not be the main thread of your program. It is advisable to keep the callback functions short and not resource-intensive. Consider strategies which use the callback only to record that a change has occurred and then act on that change later – perhaps in a separate thread, perhaps after `pend_event()` has completed.

The *cb_info* parameter supplied to the callback needs special attention, as it is the only non-Epics information passed. The *cb_info* parameter will be a tuple containing (index, self) where *index* is the key for the `callbacks` dictionary for the PV and *self* is PV object. A principle use of this tuple is to **remove the current callback** if an error happens, as for example in GUI code if the widget that the callback is meant to update disappears.

3.5 User-supplied Connection Callback functions

A *connection* callback is a user-defined function that is called when the connection status of a PV changes – that is, when a PV initially connects, disconnects or reconnects due to the process serving the PV going away, or loss of network connection. A connection callback can be specified when a PV is created, or can be added by appending to

the `connection_callbacks` list. If there is more than one connection callback defined, they will all be run when the connection state changes.

A connection callback should be prepared to receive the following keyword arguments:

- *pvname*: the name of the pv
- *conn*: the connection status

where *conn* will be either `True` or `False`, specifying whether the PV is now connected. A simple example is given below.

3.6 Put with wait, put callbacks, and put_complete

Some EPICS records take a significant amount of time to fully process, and sometimes you want to wait until the processing completes before going on. There are a few ways to accomplish this. First, one can simply wait until the processing is done:

```
import epics
p = epics.PV('XXX')
p.put(1.0, wait=True)
print 'Done'
```

This will hang until the processing of the PV completes (motor moving, etc) before printing 'Done'. You can also specify a maximum time to wait – a *timeout* (in seconds):

```
p.put(1.0, wait=True, timeout=30)
```

which will wait up to 30 seconds. For the pedantic, this timeout should not be used as an accurate clock – the actual wait time may be slightly longer.

A second method is to use the 'use_complete' option and watch for the `put_complete` attribute to become `True` after a `put()`. This is somewhat more flexible than using `wait=True` as above, because you can more carefully control how often you look for a `put()` to complete, and what to do in the interim. A simple example would be:

```
p.put(1.0, use_complete=True)
waiting = True
while waiting:
    time.sleep(0.001)
    waiting = not p.put_complete
```

An additional advantage of this approach is that you can easily wait for multiple PVs to complete with python's built-in *all* function, as with:

```
pvgroup = (epics.PV('XXX'), epics.PV('YYY'), epics.PV('ZZZ'))
newvals = (1.0, 2.0, 3.0)
for pv, val in zip(pvgroup, newvals):
    pv.put(val, use_complete=True)

waiting = True
while waiting:
    time.sleep(0.001)
    waiting = all(pv.put_complete for pv in pvgroup)
print 'All puts are done!'
```

For maximum flexibility, one can all define a *put callback*, a function to be run when the `put()` has completed. This function requires a *pvname* keyword argument, but will receive no others, unless you pass in data with the *callback_data* argument (which should be dict-like) to `put()`. A simple example would be:

```
pv = epics.PV('XXX')
def onPutComplete(pvname=None, **kws):
    print 'Put done for %s' % pvname

pv.put(1.0, callback=onPutComplete)
```

3.7 Examples

Some simple examples using PVs follow.

3.7.1 Basic Use

The simplest approach is to simply create a PV and use its `value` attribute:

```
>>> from epics import PV
>>> p1 = PV('xxx.VAL')
>>> print p1.value
1.00
>>> p1.value = 2.00
```

The `print p1.value` line automatically fetches the current PV value. The `p1.value = 2.00` line does a `put()` to set the value, causing any necessary processing over the network.

The above example is equivalent to

```
>>> from epics import PV
>>> p1 = PV('xxx.VAL')
>>> print p1.get()
1.00
>>> p1.put(value = 2.00)
```

To get a string representation of the value, you can use either

```
>>> print p1.get(as_string=True)
'1.000'
```

or, equivalently

```
>>> print p1.char_value
'1.000'
```

3.7.2 Example of using info and more properties examples

A PV has many attributes. This can be seen from its *info* paragraph:

```
>>> import epics
>>> p = epics.PV('13IDA:m3')
>>> print p.info
== 13IDA:m3 (native_double) ==
    value      = 0.2
    char_value = '0.200'
    count      = 1
    type       = double
    units      = mm
```

```

precision = 3
host       = ioc13ida.cars.aps.anl.gov:5064
access     = read/write
status     = 0
severity   = 0
timestamp  = 1274809682.967 (2010-05-25 12:48:02.967364)
upper_ctrl_limit = 5.49393415451
lower_ctrl_limit = -14.5060658455
upper_disp_limit = 5.49393415451
lower_disp_limit = -14.5060658455
upper_alarm_limit = 0.0
lower_alarm_limit = 0.0
upper_warning_limit = 0.0
lower_warning_limit = 0.0
PV is internally monitored, with 0 user-defined callbacks:
=====

```

The individual attributes can also be accessed as below. Many of these (the *control attributes*, see [Table of Control Attributes](#)) will not be filled in until either the `info` attribute is accessed or until `get_ctrlvars()` is called.

```

>>> print p.type
double
>>> print p.units, p.precision, p.lower_disp_limit
mm 3 -14.5060658455

```

3.7.3 Getting a string value

It is not uncommon to want a string representation of a PVs value, for example to show in a display window or to write to some report. For string PVs and integer PVs, this is a simple task. For floating point values, there is ambiguity how many significant digits to show. EPICS PVs all have a `precision` field. which sets how many digits after the decimal place should be described. In addition, for ENUM PVs, it would be desire able to get at the name of the ENUM state, not just its integer value.

To get the string representation of a PVs value, use either the `char_value` attribute or the `as_string=True` argument to `get()`

3.7.4 Example of put ()

To put a new value to a variable, either of these two approaches can be used:

```

>>> import epics
>>> p = epics.PV('XXX')
>>> p.put(1.0)

```

Or (equivalently):

```

>>> import epics
>>> p = epics.PV('XXX')
>>> p.value = 1.0

```

The `value` attribute is the only attribute that can be set.

3.7.5 Example of simple callback

It is often useful to get a notification of when a PV changes. In general, it would be inconvenient (and possibly inefficient) to have to continually ask if a PV's value has changed. Instead, it is better to set a *callback* function: a function to be run when the value has changed.

A simple example of this would be:

```
import epics
import time
def onChanges(pvname=None, value=None, char_value=None, **kw):
    print 'PV Changed! ', pvname, char_value, time.ctime()

mypv = epics.PV(pvname)
mypv.add_callback(onChanges)

print 'Now wait for changes'

t0 = time.time()
while time.time() - t0 < 60.0:
    time.sleep(1.e-3)
print 'Done.'
```

This first defines a *callback function* called *onChanges()* and then simply waits for changes to happen. Note that the callback function should take keyword arguments, and generally use ***kw* to catch all arguments. See [User-supplied Callback functions](#) for more details.

3.7.6 Example of connection callback

A connection callback:

```
#
# example of using a connection callback that will be called
# for any change in connection status

import epics
import time
import sys
from pvnames import motor1

write = sys.stdout.write
def onConnectionChange(pvname=None, conn=None, **kws):
    write('PV connection status changed: %s %s\n' % (pvname, repr(conn)))
    sys.stdout.flush()

def onValueChange(pvname=None, value=None, host=None, **kws):
    write('PV value changed: %s (%s) %s\n' % (pvname, host, repr(value)))
    sys.stdout.flush()
mypv = epics.PV(motor1,
                 connection_callback= onConnectionChange,
                 callback= onValueChange)

mypv.get()

write('Now waiting, watching values and connection changes:\n')
t0 = time.time()
```

```
while time.time()-t0 < 300:  
    time.sleep(0.01)
```


CA: LOW-LEVEL CHANNEL ACCESS MODULE

The `ca` module provides a low-level wrapping of the EPICS Channel Access (CA) library, using ctypes. Most users of the `epics` module will not need to be concerned with most of the details here, and will instead use the simple functional interface (`epics.caget()`, `epics.caput()` and so on), or use the `epics.PV` class to create and use `epics PV` objects.

4.1 General description, difference with C library

The goal of the `ca` module is to provide a fairly complete mapping of the C interface to the CA library while also providing a pleasant Python experience. It is expected that anyone looking into the details of this module is somewhat familiar with Channel Access and knows where to consult the [Channel Access Reference Documentation](#). This document focuses on the differences with the C interface, assuming a general understanding of what the functions are meant to do.

4.1.1 Name Mangling

As a general rule, a CA function named `ca_XXX` in the C library will have the equivalent function called `XXX` in the `ca` module. This is because the intention is that one will import the `ca` module with

```
>>> from epics import ca
```

so that the Python function `ca.XXX()` will correspond to the C function `ca_XXX`. That is, the CA library called its functions `ca_XXX` because C does not have namespaces. Python does have namespaces, and so they are used.

Similar name *un-mangling* also happens with the DBR prefixes for constants, held here in the `dbr` module. Thus, the C constant `DBR_STRING` becomes `dbr.STRING` in Python.

4.1.2 Other Changes and Omissions

Several functions in the C version of the CA library are not implemented in the Python module. Most of these unimplemented functions are currently seen as unnecessary for Python, though some of these could be added without much trouble if needed. See [Omissions](#) for further details.

In addition, while the CA library supports several DBR types in C, not all of these are supported in Python. Only native types and their `DBR_TIME` and `DBR_CTRL` variants are supported here. The `DBR_STS` and `DBR_GR` variants are not, as they are subsets of the `DBR_CTRL` type, and space optimization is not something you'll be striving for with

Python. Several `dbf_XXX` functions are also not supported, as they appear to be needed only to be able to dynamically allocate memory, which is not necessary in Python.

4.2 Initialization, Finalization, and Life-cycle

The Channel Access library must be initialized before it can be used. There are 3 main reasons for this need:

1. CA requires a context model (preemptive callbacks or non-preemptive callbacks) to be specified before any actual calls can be made.
2. the ctypes interface requires that the shared library be loaded before it is used.
3. ctypes also requires that references to the library and callback functions be kept for the life-cycle of CA-using part of a program (or else they will be garbage collected).

As far as is possible, the `ca` module hides the details of the CA lifecycle from the user, so that it is not necessary to worry about explicitly initializing a Channel Access session. Instead, the library is initialized as soon as it is needed, and intervention is really only required to change default settings. The `ca` module also handles finalizing the CA session, so that core-dumps and warning messages do not happen due to CA still being ‘alive’ as a program ends.

Because some users may wish to customize the initialization and finalization process, the detailed steps will be described here. These initialization and finalization tasks are handled in the following way:

- The `libca` variable in the `ca` module holds a permanent, global reference to the CA shared object library (DLL).
- the function `initialize_libca()` is called to initialize libca. This function takes no arguments, but does use the global Boolean `PREEMPTIVE_CALLBACK` (default value of `True`) to control whether preemptive callbacks are used.
- the function `finalize_libca()` is used to finalize libca. Normally, this function is registered to be called when a program ends with `atexit.register()`. Note that this only gets called on a graceful shutdown. If the program crashes (for a non-CA related reason, for example), this finalization may not be done, and connections to Epics Variables may not be closed completely on the Channel Access server.

`ca.PREEMPTIVE_CALLBACK`

sets whether preemptive callbacks will be used. The default value is `True`. If you wish to run without preemptive callbacks this variable *MUST* be set before any other use of the CA library. With preemptive callbacks enabled, EPICS communication will not require client code to continually poll for changes. With preemptive callback disables, you will need to frequently poll epics with `pend_io()` and `func:pend_event`.

`ca.DEFAULT_CONNECTION_TIMEOUT`

sets the default *timeout* value (in seconds) for `connect_channel()`. The default value is `2.0`

`ca.AUTOMONITOR_MAXLENGTH`

sets the default array length (ie, how many elements an array has) above which automatic conversion to numpy arrays *and* automatic monitoring for PV variables is suppressed. The default value is `65536`. To be clear: waveforms with fewer elements than this value will be automatically monitored changes, and will be converted to numpy arrays (if numpy is installed). Larger waveforms will not be automatically monitored.

advanced-arrays-label and *Strategies for working with large arrays* for more details.

4.3 Using the CA module

Many general-purpose CA functions that deal with general communication and threading contexts are very close to the C library:

`ca.initialize_libca()`
 This initializes the CA library. This must be called prior to any actual use of the CA library, but it is called automatically by the `withCA()` decorator, so you should never need to call this in a real program.

`ca.context_create()`

`ca.create_context()`
 This will create a new context, using the value of `PREEMPTIVE_CALLBACK` to set the context type. Note that CA library function has the irritating name of `context_create`. Both that and `create_context` (which is more consistent with the Verb_Object of the rest of the CA library) are allowed.

`ca.context_destroy()`

`ca.destroy_context()`
 This will destroy the current context.

`ca.current_context()`
 This returns an integer value for the current context.

`ca.attach_context(context)`
 This attaches to the context supplied.

`ca.detach_context()`
 This detaches from the current context.

`ca.use_initial_context()`
 This attaches to the context created when libca is initialized. Using this function is recommended when writing Threaded programs that using CA. See [Using Python Threads](#) for further discussion.

`ca.client_status(context, level)`
 Print (to stderr) information about Channel Access status, including status for each channel, and search and connection statistics.

`ca.version()`
 Print Channel Access version string. Currently, this should report '4.13'

`ca.message(status)`
 Print a message corresponding to a Channel Access status return value.

`ca.flush_io()`

`ca.replace_printf_handler(fcn)`
 replace the `printf()` function with the supplied function (defaults to `sys.stderr.write()`)

`ca.pend_io([t=1.0])`

`ca.pend_event([t=1.e-5])`

`ca.poll([evt=1.e-5[, iot=1.0]])`
 a convenience function which is equivalent to:

```

    pend_event(evt)
    pend_io_(iot)

```

4.3.1 Creating and Connecting to Channels

The basic channel object is the Channel ID or `chid`. With the CA library (and `ca` module), one creates and acts on the `chid` values. These are simply `ctypes.c_long` (C long integers) that hold the memory address of the C representation of the channel, but it is probably a good idea to treat these as object instances.

`ca.create_channel(pvname, [connect=False, [callback=None, auto_cb=True]])`
 creates a channel, returning the Channel ID `chid` used by other functions to identify this channel.

Parameters

- **pvname** – the name of the PV to create.
- **connect** (True/False) – whether to (try to) connect to PV as soon as possible.
- **callback** (None or callable.) – user-defined Python function to be called when the connection state changes.
- **auto_cb** (True/False) – whether to automatically use an internal callback.

The user-defined callback function should be prepared to accept keyword arguments of

- *pvname* name of PV
- *chid* `chid` Channel ID
- *conn* True/False: whether channel is connected.

If *auto_cb* is True, an internal connection callback is used so that you should not need to explicitly connect to a channel, unless you are having difficulty with dropped connections.

`ca.connect_channel(chid[, timeout=None[, verbose=False]])`
explicitly connect to a channel (usually not needed, as implicit connection will be done when needed), waiting up to timeout for a channel to connect. It returns the connection state, True or False.

Parameters

- **chid** – `chid` Channel ID
- **timeout** (float or None.) – maximum time to wait for connection.
- **verbose** – whether to print out debugging information

if *timeout* is None, the value of `DEFAULT_CONNECTION_TIMEOUT` is used (usually 2.0 seconds).

Normally, channels will connect in milliseconds, and the connection callback will succeed on the first attempt.

For un-connected Channels (that are nevertheless queried), the ‘ts’ (timestamp of last connection attempt) and ‘failures’ (number of failed connection attempts) from the `_cache` will be used to prevent spending too much time waiting for a connection that may never happen.

Many other functions that require a valid Channel ID, but not necessarily a connected Channel. These functions are essentially identical to the CA library are:

`ca.name(chid)`
return PV name for Channel.

`ca.host_name(chid)`
return host name and port serving Channel.

`ca.element_count(chid)`
return number of elements in Channel’s data.

`ca.read_access(chid)`
return *read access* for a Channel: 1 for True, 0 for False.

`ca.write_access(chid)`
return *write access* for a channel: 1 for True, 0 for False.

`ca.field_type(chid)`
return the integer DBR field type. See the *ftype* column from *Table of DBR Types*.

`ca.clear_channel(chid)`
clear the channel.

`ca.state(chid)`
return the state of the channel.

A few additional pythonic functions have been added:

`ca.isConnected(chid)`
returns `dbf.CS_CONN==state(chid)` ie `True` for a connected channel or `False` for an unconnected channel.

`ca.access(chid)`
returns a string describing read/write access: one of *no access*, *read-only*, *write-only*, or *read/write*

`ca.promote_type(chid[, use_time=False[, use_ctrl=False]])`
promotes the native field type of a `chid` to its `TIME` or `CTRL` variant. See [Table of DBR Types](#). Returns the integer corresponding to the promoted field value.

`ca._cache`
The `ca` module keeps a global cache of Channels that holds connection status and a bit of internal information for all known PVs. This cache is not intended for general use.

`ca.show_cache([print_out=True])`
this function will print out a listing of PVs in the current session to standard output. Use the `print_out=False` option to be returned the listing instead of having it printed.

4.3.2 Interacting with Connected Channels

Once a `chid` is created and connected there are several ways to communicating with it. These are primarily encapsulated in the functions `get()`, `put()`, and `create_subscription()`, with a few additional functions for retrieving specific information.

These functions are where this python module differs the most from the underlying CA library, and this is mostly due to the underlying CA function requiring the user to supply `DBR TYPE` and count as well as `chid` and allocated space for the data. In python none of these is needed, and keyword arguments can be used to specify such options.

`ca.get(chid[, ftype=None[, count=None[, as_string=False[, as_numpy=True[, wait=True[, timeout=None]]]]])`
return the current value for a Channel. Note that there is not a separate form for array data.

Parameters

- **chid** (`ctypes.c_long`) – `chid` Channel ID
- **ftype** (integer or `None`) – field type to use (native type is default)
- **count** (integer or `None`) – maximum element count to return (full data returned by default)
- **as_string** (`True/False`) – whether to return the string representation of the value. See notes below.
- **as_numpy** (`True/False`) – whether to return the Numerical Python representation for array / waveform data.
- **wait** (`True/False`) – whether to wait for the data to be received, or return immediately.
- **timeout** (float or `None`) – maximum time to wait for data before returning `None`.

`get()` returns the value for the PV with channel ID `chid` or `None`, which indicates an *incomplete get*

For a listing of values of `ftype`, see [Table of DBR Types](#). The optional `count` can be used to limit the amount of data returned for array data from waveform records.

The `as_string` option warrants special attention: The feature is not as complete as as the `as_string` argument for `PV.get()`. Here, a string representing the value will always be returned. For Enum types, the name of the Enum state will be returned. For waveforms of type `CHAR`, the string representation will be returned. For other

waveforms (with *count* > 1), a string like `<array count=3, type=1>` will be returned. For all other types the result will from Python's `str()` function.

The *as_numpy* option will cause an array value to be returned as a numpy array. This is only applied if numpy can be imported. See [Strategies for working with large arrays](#) for a discussion of strategies for how to best deal with very large arrays.

The *wait* option controls whether to wait for the data to be received over the network and actually return the value, or to return immediately after asking for it to be sent. If *wait=False* (that is, immediate return), the *get* operation is said to be *incomplete*. The data will be still be received (unless the channel is disconnected) eventually but stored internally, and can be read later with `get_complete()`. Using *wait=False* can be useful in some circumstances. See [Strategies for connecting to a large number of PVs](#) for a discussion.

The *timeout* option sets the maximum time to wait for the data to be received over the network before returning `None`. Such a timeout could imply that the channel is disconnected or that the data size is larger or network slower than normal. In that case, the *get* operation is said to be *incomplete*, and the data may become available later with `get_complete()`.

See [The wait and timeout options for get\(\), ca.get_complete\(\)](#) for further discussion of the *wait* and *timeout* options and the associated `get_complete()` function.

```
ca.get_complete(chid[, ftype=None[, count=None[, as_string=False[, as_numpy=True[, time-  
out=None]]]]])
```

return the current value for a Channel, completing an earlier incomplete `get()` that returned `None`, either because *wait=False* was used or because the data transfer did not complete before the timeout passed.

Parameters

- **chid** (*ctypes.c_long*) – chid Channel ID
- **ftype** (*integer*) – field type to use (native type is default)
- **count** (*integer*) – maximum element count to return (full data returned by default)
- **as_string** (True/False) – whether to return the string representation of the value. See notes below.
- **as_numpy** (True/False) – whether to return the Numerical Python representation for array / waveform data.
- **timeout** (float or None) – maximum time to wait for data before returning `None`.

This function will return `None` if the previous `get()` actually completed, or if this data transfer also times out. See [The wait and timeout options for get\(\), ca.get_complete\(\)](#) for further discussion.

```
ca.put(chid, value[, wait=False[, timeout=30[, callback=None[, callback_data=None]]]])
```

sets the Channel to a value, with options to either wait (block) for the process to complete, or to execute a supplied callback function when the process has completed. The chid and value are required.

Parameters

- **chid** (*ctypes.c_long*) – chid Channel ID
- **wait** (True/False) – whether to wait for processing to complete (or time-out) before returning.
- **timeout** (float or None) – maximum time to wait for processing to complete before returning anyway.
- **callback** (None or callable) – user-supplied function to run when processing has completed.
- **callback_data** – extra data to pass on to a user-supplied callback function.

`put()` returns 1 on success and -1 on timed-out

Specifying a callback will override setting `wait=True`. This callback function will be called with keyword arguments

`pvname=pvname, data=callback_data`

For more on this *put* callback, see [User-supplied Callback functions](#) below.

```
ca.create_subscription(chid[, use_time=False[, use_ctrl=False[, mask=None[, callback=None]])
```

create a *subscription to changes*, The user-supplied callback function will be called on any changes to the PV.

Parameters

- **use_time** (True/False) – whether to use the TIME variant for the PV type
- **use_ctrl** (True/False) – whether to use the CTRL variant for the PV type
- **mask** (*integer*) – bitmask (combination of `dbr.DBE_ALARM`, `dbr.DBE_LOG`, `dbr.DBE_VALUE`) to control which changes result in a callback. Defaults to `DEFAULT_SUBSCRIPTION_MASK`.
- **callback** (None or callable) – user-supplied callback function

Return type tuple containing (*callback_ref*, *user_arg_ref*, *event_id*)

The returned tuple contains *callback_ref* and *user_arg_ref* which are references that should be kept for as long as the subscription lives (otherwise they may be garbage collected, causing no end of trouble). *event_id* is the id for the event (useful for clearing a subscription).

For more on writing the user-supplied callback, see [User-supplied Callback functions](#) below.

Warning: *event_id* is the id for the event (useful for clearing a subscription). You **must** keep the returned tuple in active variables, either as a global variable or as data in an encompassing class. If you do *not* keep this data, the return value will be garbage collected, the C-level reference to the callback will disappear, and you will see coredumps.

On Linux, a message like:

```
python: Objects/funcobject.c:451: func_dealloc: Assertion 'g->gc.gc_refs != (-2)' failed.
Abort (core dumped)
```

is a hint that you have *not* kept this data.

`ca.DEFAULT_SUBSCRIPTION_MASK`

This value is the default subscription type used when calling `create_subscription()` with `mask=None`. It is also used by default when creating a PV object with `auto_monitor` is set to `True`.

The initial default value is `dbr.DBE_ALARM|dbr.DBE_VALUE` (i.e. update on alarm changes or value changes which exceeds the monitor deadband.) The other possible flag in the bitmask is `dbr.DBE_LOG` for archive-deadband changes.

If this value is changed, it will change the default for all subsequent calls to `create_subscription()`, but it will not change any existing subscriptions.

`ca.clear_subscription(event_id)`

clears a subscription given its *event_id*.

Several other functions are provided:

`ca.get_timestamp(chid)`

return the timestamp of a channel – the time of last update.

`ca.get_severity(chid)`
return the severity of a channel.

`ca.get_precision(chid)`
return the precision of a channel. For channels with native type other than FLOAT or DOUBLE, this will be 0.

`ca.get_enum_strings(chid)`
return the list of names for ENUM states of a Channel. Returns `None` for non-ENUM Channels.

`ca.get_ctrlvars(chid)`
returns a dictionary of CTRL fields for a Channel. Depending on the native data type, the keys in this dictionary may include *Table of Control Attributes*

Table of Control Attributes

<i>attribute</i>	<i>data types</i>
status	
severity	
precision	0 for all but double, float
units	
enum_strs	enum only
upper_disp_limit	
lower_disp_limit	
upper_alarm_limit	
lower_alarm_limit	
upper_warning_limit	
lower_warning_limit	
upper_ctrl_limit	
lower_ctrl_limit	

Note that *enum_strs* will be a tuple of strings for the names of ENUM states.

`ca.get_timevars(chid)`
returns a dictionary of TIME fields for a Channel. This will contain a *status*, *severity*, and *timestamp* key.

4.3.3 Synchronous Groups

Synchronous Groups can be used to ensure that a set of Channel Access calls all happen together, as if in a *transaction*. Synchronous Groups work in PyEpics as of version 3.0.10, but more testing is probably needed.

The idea is to first create a synchronous group, then add a series of `sg_put()` and `sg_get()` which do not happen immediately, and finally block while all the channel access communication is done for the group as a unit. It is important to *not* issue `pend_io()` during the building of a synchronous group, as this will cause pending `sg_put()` and `sg_get()` to execute.

`ca.sg_create()`
create synchronous group. Returns a *group id*, *gid*, which is used to identify this group and is passed to all other synchronous group commands.

`ca.sg_delete(gid)`
delete a synchronous group

`ca.sg_block(gid[, t=10.0])`
block for a synchronous group to complete processing

`ca.sg_get(gid, chid[, ftype=None[, as_string=False[, as_numpy=True]]])`
perform a *get* within a synchronous group.

This function will not immediately return the value, of course, but the address of the underlying data.

After the `sg_block()` has completed, you must use `_unpack()` to convert this data address to the actual value(s).

See example below.

`ca.sg_put(gid, chid, value)`
perform a *put* within a synchronous group. This *put* cannot wait for completion.

`ca.sg_test(gid)`
test whether a synchronous group has completed.

`ca.sg_reset(gid)`
resets a synchronous group

An example use of a synchronous group:

```
from epics import ca
import time

pvs = ('X1.VAL', 'X2.VAL', 'X3.VAL')
chids = [ca.create_channel(pvname) for pvname in pvs]

for chid in chids:
    ca.connect_channel(chid)
    ca.put(chid, 0)

# create synchronous group
sg = ca.sg_create()

# get data pointers from ca.sg_get
data = [ca.sg_get(sg, chid) for chid in chids]

print 'Now change these PVs for the next 10 seconds'
time.sleep(10.0)

print 'will now block for i/o'
ca.sg_block(sg)
#
# CALL ca._unpack with data points and chid to extract data
for pvname, dat, chid in zip(pvs, data, chids):
    val = ca._unpack(dat, chid=chid)
    print "%s = %s" % (pvname, str(val))

ca.sg_reset(sg)

# Now a SG Put
print 'OK, now we will put everything back to 0 synchronously'

for chid in chids:
    ca.sg_put(sg, chid, 0)

print 'sg_put done, but not blocked / committed. Sleep for 5 seconds '
time.sleep(5.0)
ca.sg_block(sg)
print 'done.'
```

4.4 Implementation details

The details given here should mostly be of interest to those looking at the implementation of the `ca` module, those interested in the internals, or those looking to translate lower-level C or Python code to this module.

4.4.1 DBR data types

Table of DBR Types

<i>CA type</i>	<i>integer ftype</i>	<i>Python ctypes type</i>
string	0	string
int	1	integer
short	1	integer
float	2	double
enum	3	integer
char	4	byte
long	5	integer
double	6	double
time_string	14	
time_int	15	
time_short	15	
time_float	16	
time_enum	17	
time_char	18	
time_long	19	
time_double	20	
ctrl_string	28	
ctrl_int	29	
ctrl_short	29	
ctrl_float	30	
ctrl_enum	31	
ctrl_char	32	
ctrl_long	33	
ctrl_double	34	

4.4.2 PySEVCHK and ChannelAccessException: checking CA return codes

exception `ca.ChannelAccessException`

This exception is raised when the `ca` module experiences unexpected behavior and must raise an exception

`ca.PySEVCHK(func_name, status[, expected=dbr.ECA_NORMAL])`

This checks the return *status* returned from a `libca.ca_***` and raises a `ChannelAccessException` if the value does not match the *expected* value.

The message from the exception will include the *func_name* (name of the Python function) and the CA message from message.

`ca.withSEVCHK()`

this decorator handles the common case of running `PySEVCHK()` for a function whose return value is from a `libca.ca_***` function and whose return value should be `dbr.ECA_NORMAL`.

4.4.3 Function Decorators

In addition to `withSEVCHK()`, several other decorator functions are used heavily inside of `ca.py` or are available for your convenience.

`ca.withCA()`

ensures that the CA library is initialized before many CA functions are called. This prevents, for example, one creating a channel ID before CA has been initialized.

`ca.withCHID()`

ensures that CA functions which require a `chid` as the first argument actually have a `chid` as the first argument. This is not a highly robust test (it actually checks for a `ctypes.c_long` or `int`) but is useful enough to catch most errors before they would cause a crash of the CA library.

`ca.withConnectedCHID()`

ensures that the first argument of a function is a connected `chid`. This test is (intended to be) robust, and will (try to) make sure a `chid` is actually connected before calling the decorated function.

`ca.withInitialContext()`

ensures that the called function uses the threading context initially defined. The See [Using Python Threads](#) for further discussion.

4.4.4 Unpacking Data from Callbacks

Throughout the implementation, there are several places where data returned by the underlying CA library needs to be converted to Python data. This is encapsulated in the `_unpack()` function. In general, you will not have to run this code, but there is one exception: when using `sg_get()`, the values returned will have to be unpacked with this function.

`ca._unpack(cdata, chid=None[, count=None[, ftype=None[, as_numpy=None]]])`

This takes the ctypes data `cdata` and returns the Python data.

Parameters

- **cdata** – cdata as returned by internal libca functions, and `sg_get()`.
- **chid** – channel ID (optional: used for determining count and ftype)
- **count** – number of elements to fetch (defaults to element count of `chid` or 1)
- **ftype** – data type of channel (defaults to native type of `chid`)
- **as_numpy** (True/False) – whether to convert to numpy array.

4.5 User-supplied Callback functions

User-supplied callback functions can be provided for both `put()` and `create_subscription()`. Note that callbacks for *PV* objects are slightly different: see [User-supplied Callback functions](#) in the `pv` module for details.

When defining a callback function to be run either when a `put()` completes or on changes to the Channel, as set from `create_subscription()`, it is important to know two things:

1. how your function will be called.
2. what is permissible to do inside your callback function.

In both cases, callbacks will be called with keyword arguments. You should be prepared to have them passed to your function. Use `**kw` unless you are very sure of what will be sent.

For callbacks sent when a `put ()` completes, your function will be passed these:

- *pvname* : the name of the pv
- *data*: the user-supplied `callback_data` (defaulting to `None`).

For subscription callbacks, your function will be called with keyword/value pairs that will include:

- *pvname*: the name of the pv
- *value*: the latest value
- *count*: the number of data elements
- *ftype*: the numerical CA type indicating the data type
- *status*: the status of the PV (1 for OK)
- *chid*: the integer address for the channel ID.

Depending on the data type, and whether the CTRL or TIME variant was used, the callback function may also include some of these as keyword arguments:

- *enum_strs*: the list of enumeration strings
- *precision*: number of decimal places of precision.
- *units*: string for PV units
- *severity*: PV severity
- *timestamp*: timestamp from CA server.

Note that a the user-supplied callback will be run *inside* a CA function, and cannot reliably make any other CA calls. It is helpful to think “this all happens inside of a `pend_event ()` call”, and in an epics thread that may or may not be the main thread of your program. It is advisable to keep the callback functions short and not resource-intensive. Consider strategies which use the callback only to record that a change has occurred and then act on that change later – perhaps in a separate thread, perhaps after `pend_event ()` has completed.

4.6 Omissions

Several parts of the CA library are not implemented in the Python module. These are currently seen as unneeded (with notes where appropriate for alternatives), though they could be added on request.

`ca.ca_add_exception_event ()`

Not implemented: Python exceptions are raised where appropriate and can be used in user code.

`ca.ca_add_fd_registration ()`

Not implemented

`ca.ca_replace_access_rights_event ()`

Not implemented

`ca.ca_client_status ()`

Not implemented

`ca.ca_set_puser ()`

Not implemented : it is easy to pass user-defined data to callbacks as needed.

`ca.ca_puser ()`

Not implemented: it is easy to pass user-defined data to callbacks as needed.

`ca.ca_SEVCHK()`

Not implemented: the Python function `PySEVCHK()` is approximately the same.

`ca.ca_signal()`

Not implemented: the Python function `PySEVCHK()` is approximately the same.

`ca.ca_test_event()`

Not implemented: this appears to be a function for debugging events. These are easy enough to simulate by directly calling Python callback functions.

`ca.ca_dump_dbr()`

Not implemented

In addition, not all *DBR* types in the CA C library are supported.

Only native types and their *DBR_TIME* and *DBR_CTRL* variants are supported: *DBR_STS* and *DBR_GR* variants are not. Several *dbr_XXX* functions are also not supported, as they are needed only to dynamically allocate memory.

4.7 CATHread class

`class ca.CATHread(group=None[, target=None[, name=None[, args=()[, kwargs={}]]]])`

create a CA-aware subclass of a standard Python `threading.Thread`. See the standard library documentation for further information on how to use Thread objects.

A *CATHread* simply runs `use_initial_context()` prior to running each target function, so that `use_initial_context()` does not have to be explicitly put inside the target function.

The See *Using Python Threads* for further discussion.

4.8 Examples

Here are some example sessions using the `ca` module.

4.8.1 Create, Connect, Get Value of Channel

Note here that several things have been simplified compare to using CA in C: initialization and creating a main-thread context are handled, and connection of channels is handled in the background:

```
from epics import ca
chid = ca.create_channel('XXX:m1.VAL')
count = ca.element_count(chid)
ftype = ca.field_type(chid)
print "Channel ", chid, count, ftype
value = ca.get()
print value
```

4.8.2 Put, waiting for completion

Here we set a PVs value, waiting for it to complete:

```
from epics import ca
chid = ca.create_channel('XXX:m1.VAL')
ca.put(chid, 1.0, wait=True)
```

The `put()` method will wait to return until the processing is complete.

4.8.3 Define a callback to Subscribe to Changes

Here, we *subscribe to changes* for a PV, which is to say we define a callback function to be called whenever the PV value changes. In the case below, the function to be called will simply write the latest value out to standard output:

```
from epics import ca
import time
import sys

# define a callback function. Note that this should
# expect certain keyword arguments, including 'pvname' and 'value'
def onChanges(pvname=None, value=None, **kw):
    fmt = 'New Value: %s value=%s, kw=%s\n'
    sys.stdout.write(fmt % (pvname, str(value), repr(kw)))
    sys.stdout.flush()

# create the channel
mypv = 'XXX.VAL'
chid = ca.create_channel(mypv)

# subscribe to events giving a callback function
eventID = ca.create_subscription(chid, callback=onChanges)

# now we simply wait for changes
t0 = time.time()
while time.time()-t0 < 10.0:
    time.sleep(0.001)
```

It is **vital** that the return value from `create_subscription()` is kept in a variable so that it cannot be garbage collected. Failure to keep this value will cause trouble, including almost immediate segmentation faults (on Windows) or seemingly inexplicable crashes later (on linux).

4.8.4 Define a connection callback

Here, we define a connection callback – a function to be called when the connection status of the PV changes. Note that this will be called on initial connection:

```
import epics
import time

def onConnectionChange(pvname=None, conn=None, chid=None):
    print 'ca connection status changed: ', pvname, conn, chid

# create channel, provide connection callback
motor1 = '13IDC:m1'
chid = epics.ca.create_channel(motor1, callback=onConnectionChange)

print 'Now waiting, watching values and connection changes:'
t0 = time.time()
while time.time()-t0 < 30:
    time.sleep(0.001)
```

This will run the supplied callback soon after the channel has been created, when a successful connection has been made. Note that the callback should be prepared to accept keyword arguments of *pvname*, *chid*, and *conn* for the PV

name, channel ID, and connection state (`True` or `False`).

DEVICES: COLLECTIONS OF PVS

5.1 Overview

The `device` module provides a simple interface to a collection of PVs. Here an `epics.device.Device` is an object holding a set of PVs, all sharing a prefix, but having many *attributes*. Many PVs will have names made up of *prefix+attribute*, with a common prefix for several related PVs. This almost describes an Epics Record, but as it is concerned only with PV names, the mapping to an Epics Record is not exact. On the other hand, the concept of a *device* is more flexible than a predefined Epics Record as it can actually hold PVs from several different records.:

```
motor1 = epics.Device('XXX:motor1.', attr=('VAL', 'RBV', 'DESC', 'RVAL',
                                           'LVIO', 'HLS', 'LLS'))

motor1.put('VAL', 1)
print 'Motor %s = %f' % ( motor1.get('DESC'), motor1.get('RBV'))

motor1.VAL 0
print 'Motor %s = %f' % ( motor1.DESC, motor1.RBV )
```

While useful on its own like this, the real point of a *device* is as a base class, to be inherited and extended. In fact, there is a more sophisticated Motor device described below at [Epics Motor Device](#)

```
class device.Device (prefix=None[, delim=' ', attrs=None ])
```

The attribute PVs are built as needed and held in an internal buffer `self._pvs`. This class is kept intentionally simple so that it may be subclassed.

To pre-load attribute names on initialization, provide a list or tuple of attributes with the *attr* option.

Note that *prefix* is actually optional. When left off, this class can be used as an arbitrary container of PVs, or to turn any subclass into an `epics.Device`.

In general, PV names will be mapped as `prefix+delim+attr`. See `add_pv()` for details of how to override this.

```
device.PV (attr[, connect=True[, **kw]])
    returns the PV object for a device attribute. The connect argument and any other keyword arguments are passed to epics.PV().

device.put (attr, value[, wait=False[, timeout=10.0 ]])
    put an attribute value, optionally wait for completion or up to a supplied timeout value

device.get (attr[, as_string=False ])
    get an attribute value, option as_string returns a string representation

device.add_callback (attr, callback)
    add a callback function to an attribute PV, so that the callback function will be run when the attribute's value changes
```

`device.add_pv (pyname[, attr=None[, **kw]])`

adds an explicitly names `epics.PV()` to the device even though it may violate the normal naming rules (in which `attr` is mapped to `epics.PV(prefix+delim+attr)`. That is, one can say:

```
import epics
m1 = epics.Device('XXX:m1', delim='.')
m1.add_pv('XXX:m2.VAL', attr='other')
print m1.VAL      # print value of XXX:m1.VAL
print m1.other    # prints value of XXX:m2.VAL
```

`device.save_state()`

return a dictionary of all current values – the “current state”.

`device.restore_state (state)`

restores a saved state, as saved with `save_state()`

`device.write_state (fname[, state=None])`

write a saved state to a file. If no state is provide, the current state is written.

`device.read_state (fname[, restore=False])`

reads a state from a file, as written with `write_state()`, and returns it. If “restore” is `True`, the read state will be restored.

`device._pvs`

a dictionary of PVs making up the device.

5.2 Epics Motor Device

The Epics Motor record has over 100 fields associated with it. Of course, it is often preferable to think of 1 Motor with many attributes than 100 or so separate PVs. Many of the fields of the Motor record are interrelated and influence other settings, including limits on the range of motion which need to be respected, and which may send notifications when they are violated. Thus, there is a fair amount of functionality for a Motor. Typically, the user just wants to move the motor by setting its drive position, but a fully enabled Motor should allow the use to change and read many of the Motor parameters.

The `Motor` class helps the user create and use Epics motors. A simple example use would be:

```
import epics
m1 = epics.Motor('XXX:m1')

print 'Motor: ', m1.DESC , ' Currently at ', m1.RBV

m1.tweak_val = 0.10
m1.move(0.0, dial=True, wait=True)

for i in range(10):
    m1.tweak(dir='forward', wait=True)
    time.sleep(1.0)
    print 'Motor: ', m1.DESC , ' Currently at ', m1.RBV
```

Which will step the motor through a set of positions. You’ll notice a few features for `Motor`:

1. Motors can use English-name aliases for attributes for fields of the motor record. Thus ‘VAL’ can be spelled ‘drive’ and ‘DESC’ can be ‘description’. The Table
2. The methods for setting positions can use the User, Dial, or Step coordinate system, and can wait for completion.

5.2.1 The `epics.Motor` class

`class motor.Motor(pvname[, timeout=30.]`
 create a Motor object for a named Epics Process Variable.

Parameters

- **pvname** (*string*) – prefix name (no ‘.VAL’ needed!) of Epics Process Variable for a Motor
- **timeout** (*float*) – time (in seconds) to wait before giving up trying to connect.

Once created, a Motor should be ready to use.

```
>>> from epics import Motor
>>> m = Motor('XX:m1')
>>> print m.drive, m.description, m.slew_speed
1.030 Fine X 5.0
>>> print m.get('device_type', as_string=True)
'asynMotor'
```

A Motor has very many fields. Only a few of them are created on initialization – the rest are retrieved as needed. The motor fields can be retrieved either with an attribute or with the `get()` method. A full list of Motor attributes and their aliases for the motor record is given in *Table of Motor Attributes*.

Table of Aliases for attributes for the `epics.Motor` class, and the corresponding attribute name of the Motor Record field.

alias	Motor Record field	alias	Motor Record field
disabled	_able.VAL	moving	MOVN
acceleration	ACCL	resolution	MRES
back_accel	BACC	motor_status	MSTA
backlash	BDST	offset	OFF
back_speed	BVEL	output_mode	OMSL
card	CARD	output	OUT
dial_high_limit	DHLM	prop_gain	PCOF
direction	DIR	precision	PREC
dial_low_limit	DLLM	readback	RBV
settle_time	DLY	retry_max	RTRY
done_moving	DMOV	retry_count	RCNT
dial_readback	DRBV	retry_deadband	RDBD
description	DESC	dial_difference	RDIF
dial_drive	DVAL	raw_encoder_pos	REP
units	EGU	raw_high_limit	RHLS
encoder_step	ERES	raw_low_limit	RLLS
freeze_offset	FOFF	relative_value	RLV
move_fraction	FRAC	raw_motor_pos	RMP
hi_severity	HHSV	raw_readback	RRBV
hi_alarm	HIGH	readback_res	RRES
hihi_alarm	HIHI	raw_drive	RVAL
high_limit	HLM	dial_speed	RVEL
high_limit_set	HLS	s_speed	S
hw_limit	HLSV	s_back_speed	SBAK
home_forward	HOMF	s_base_speed	SBAS
home_reverse	HOMR	s_max_speed	SMAX
high_op_range	HOPR	set	SET

Continued on next page

Table 5.1 – continued from previous page

alias	Motor Record field	alias	Motor Record field
high_severity	HSV	stop_go	SPMG
integral_gain	ICOF	s_revolutions	SREV
jog_accel	JAR	stop	STOP
jog_forward	JOGF	t_direction	TDIR
jog_reverse	JOGR	tweak_forward	TWF
jog_speed	JVEL	tweak_reverse	TWR
last_dial_val	LDVL	tweak_val	TWV
low_limit	LLM	use_encoder	UEIP
low_limit_set	LLS	u_revolutions	UREV
lo_severity	LLSV	use_rdbl	URIP
lolo_alarm	LOLO	drive	VAL
low_op_range	LOPR	base_speed	VBAS
low_alarm	LOW	slew_speed	VELO
last_rel_val	LRLV	version	VERS
last_dial_drive	LRVL	max_speed	VMAX
last_SPMG	LSPG	use_home	ATHM
low_severity	LSV	deriv_gain	DCOF

5.2.2 methods for `epics.Motor`

`motor.get(attr[, as_string=False])`
sets a field attribute for the motor.

Parameters

- **attr** (string (from table above)) – attribute name
- **as_string** (True/ False) – whether to return string value.

Note that `get()` can return the string value, while fetching the attribute cannot do so:

```
>>> m = epics.Motor('XXX:m1')
>>> print m.device_type
0
>>> print m.get('device_type', as_string=True)
'asynMotor'
```

`motor.put(attr, value[, wait=False[, timeout=30]])`
sets a field attribute for the motor.

Parameters

- **attr** (string (from table above)) – attribute name
- **value** – value for attribute
- **wait** (True/False) – whether to wait for completion.
- **timeout** (float) – time (in seconds) to wait before giving up trying to connect.

`motor.check_limits()`
checks whether the current motor position is causing a motor limit violation, and raises a `MotorLimitException` if it is.

returns `None` if there is no limit violation.

`motor.within_limits (value[, limits='user'])`
checks whether a target value **would be** a limit violation.

Parameters

- **value** – target value
- **limits** (*string*) – one of ‘user’, ‘dial’, or ‘raw’ for which limits to consider

Return type True/False

`motor.move (val=None[, relative=None[, wait=False[, timeout=300.0[, dial=False[, raw=False[, ignore_limits=False]]]]])`
moves motor drive to position

Parameters

- **val** – value to move to (float) [Must be provided]
- **relative** – move relative to current position (T/F) [F]
- **wait** – whether to wait for move to complete (T/F) [F]
- **dial** – use dial coordinates (T/F) [F]
- **raw** – use raw coordinates (T/F) [F]
- **ignore_limits** – try move without regard to limits (T/F) [F]
- **timeout** – max time for move to complete (in seconds) [300]

Return type see below

Return codes:

None : unable to move, invalid value given -1 : target value outside limits – no move attempted -2 : with *wait=True*, wait time exceeded timeout 0 : move executed successfully

will raise an exception if a motor limit is met.

`motor.tweak (dir='forward'[, wait=False[, timeout=300.]])`
move the motor by the current *tweak value*

Parameters

- **dir** (*string*: ‘forward’ (default) or ‘reverse’) – direction of motion
- **wait** (True/False) – whether to wait for completion
- **timeout** (*float*) – max time for move to complete (in seconds) [default=300]

`motor.get_position (readback=False[, dial=False[, raw=False]])`
Returns the motor position in user, dial or raw coordinates.

Parameters

- **readback** – whether to return the readback position in the desired coordinate system. The default is to return the drive position of the motor.
- **dial** – whether to return the position in dial coordinates. The default is user coordinates.
- **raw** – whether to return the raw position. The default is user coordinates.

The “raw” and “dial” keywords are mutually exclusive. The “readback” keyword can be used in user, dial or raw coordinates.

`motor.set_position (position[dial=False[, raw=False]])`
set (that is, redefine) the current position to supplied value.

Parameters

- **position** – The new motor position
- **dial** – whether to set in dial coordinates. The default is user coordinates.
- **raw** – whether to set in raw coordinates. The default is user coordinates.

The 'raw' and 'dial' keywords are mutually exclusive.

```
motor.get_pv(attr)
    returns the PV for the corresponding attribute.

motor.set_callback(attr='drive', callback=None[, kw=None])
    sets a callback on the PV for a particular attribute.

motor.clear_callback(attr='drive')
    clears a callback on the PV for a particular attribute.

motor.show_info()
    prints out a table of attributes and their current values.
```

5.3 Other Device Examples

As defined here, an epics device provides a general way to group together a set of PVs. The examples below show how to build on this generality, and may inspire you to build your own device classes.

5.3.1 Device without a prefix

Here, we define a very simple device that does not even define a prefix. This is not much more than a collection of PVs. All PVs in the device must be *fully qualified*, as they need do not share a common prefix:

```
from epics import Device
dev = Device()
p1 = dev.PV('13IDC:m1.VAL')
p2 = dev.PV('13IDC:m2.VAL')
dev.put('13IDC:m1.VAL', 2.8)
dev.put('13IDC:m2.VAL', 3.0)
print dev.PV('13IDC:m3.DIR').get(as_string=True)
```

Note that this device cannot use the attributes based on field names.

This may not look very interesting until you consider *Device* to be a starting point for building more complicated objects by adding specialized methods.

5.3.2 Epics ai record as Device

Here is a slightly more useful example: An Epics ai (analog input record) implemented as a Device.

```
#!/usr/bin/python
"""Epics analog input record"""
import epics

class ai(epics.Device):
    "Simple analog input device"

    attrs = ('VAL', 'EGU', 'HOPR', 'LOPR', 'PREC', 'NAME', 'DESC',
```

```

        'DTYP', 'INP', 'LINR', 'RVAL', 'ROFF', 'EGUF', 'EGUL',
        'AOFF', 'ASLO', 'ESLO', 'EOFF', 'SMOO', 'HIHI', 'LOLO',
        'HIGH', 'LOW', 'HHSV', 'LLSV', 'HSV', 'LSV', 'HYST')

    def __init__(self, prefix):
        if prefix.endswith('.'):
            prefix = prefix[:-1]
        epics.Device.__init__(self, prefix, delim='.',
                               attrs=self.attrs)

```

Note that we pre-define the fields that are the *suffixes* of an Epics ai input record, and simply subclass Device with these fields. This ai class can then be used simply and cleanly as:

```

This_ai = ai('XXX.PRES')
print 'Value: ', This_ai.VAL
print 'Units: ', This_ai.EGU

```

Several of the other standard Epics records can easily be exposed as Devices in this way.

5.3.3 Epics Scaler Record as Device

And finally a slightly more complicated example: an incomplete, but very useful mapping of the Scaler Record from synApps, including methods for changing modes, and reading and writing data.

```

#!/usr/bin/python
"""Epics Scaler"""
import epics

class Scaler(epics.Device):
    """
    Simple implementation of SynApps Scaler Record.
    """
    attrs = ('CNT', 'CONT', 'TP', 'T')
    attr_kws = {'calc_enable': '%s_calcEnable.VAL'}
    chan_attrs = ('NM%i', 'S%i')
    calc_attrs = {'calc%i': '%s_calc%i.VAL', 'expr%i': '%s_calc%i.CALC'}

    _fields = ('_prefix', '_pvs', '_delim', '_nchan', '_chans')

    def __init__(self, prefix, nchan=8):
        self._nchan = nchan
        self._chans = range(1, nchan+1)

        attrs = list(self.attrs)
        for i in self._chans:
            for att in self.chan_attrs:
                attrs.append(att % i)

        epics.Device.__init__(self, prefix, delim='.', attrs=attrs)

        for key, val in self.attr_kws.items():
            self.add_pv(val % prefix, attr= key)

        for i in self._chans:
            for key, val in self.calc_attrs.items():
                self.add_pv(val % (prefix, i), attr = key % i)

```

```
def AutoCountMode(self):
    "set to autocount mode"
    self.put('CONT', 1)

def OneShotMode(self):
    "set to one shot mode"
    self.put('CONT', 0)

def CountTime(self, ctime):
    "set count time"
    self.put('TP', ctime)

def Count(self, ctime=None):
    "set count, with optional counttime"
    if ctime is not None:
        self.CountTime(ctime)
    self.put('CNT', 1)

def EnableCalcs(self):
    "enable calculations"
    self.put('calc_enable', 1)

def setCalc(self, i, calc):
    "set the calculation for scaler i"
    attr = 'expr%i' % i
    self.put(attr, calc)

def getNames(self):
    "get all names"
    return [self.get('NM%i' % i) for i in self._chans]

def Read(self, use_calcs=False):
    "read all values"
    attr = 'S%i'
    if use_calcs:
        attr = 'calc%i'
    return [self.get(attr % i) for i in self._chans]
```

Note that we can then create a scaler object from its base PV prefix, and use methods like `Count()` and `Read()` without directly invoking epics calls:

```
s1 = Scaler('XXX:scaler1')
s1.setCalc(2, '(B-2000*A/10000000.)')
s1.enableCalcs()
s1.OneShotMode()
s1.Count(t=5.0)
print 'Names: ', s1.getNames()
print 'Raw values: ', s1.Read(use_calcs=False)
print 'Calc values: ', s1.Read(use_calcs=True)
```


ALARMS: RESPOND WHEN A PV GOES OUT OF RANGE

6.1 Overview

The `alarm` module provides an `Alarm` object to specify an alarm condition and what to do when that condition is met.

```
class alarm.Alarm(pvname[, comparison=None[, trip_point=None[, callback=None[, alert_delay=10 ] ] ])
```

creates an alarm object.

param pvname name of Epics PV (string)

param comparison operation used to compare PV value to trip_point.

type comparison string or callable. Built in comparisons are listed in *Table of Alarm Operators*.

param trip_point value that will trigger the alarm

param callback user-defined callback function to be run when the PVs value meets the alarm condition

type callback callable or None

param alert_delay time (in seconds) to wait before executing another alarm callback.

The alarm works by checking the value of the PV each time it changes. If the new value is outside the acceptable range (violates the trip point), then the user-supplied callback function is run. This callback could be set to send a message or to take some other course of action.

The comparison supplied can either be a string as listed in *Table of Alarm Operators* or a custom callable function which takes the two values (PV.value, trip_point) and returns `True` or `False` based on those values.

Table of built-in Operators for Alarms:

<i>operator</i>	Python operator
'eq', '=='	<code>__eq__</code>
'ne', '!='	<code>__ne__</code>
'le', '<='	<code>__le__</code>
'lt', '<'	<code>__lt__</code>
'ge', '>='	<code>__ge__</code>
'gt', '>'	<code>__gt__</code>

The `alert_delay` prevents the alarm callback from being called too many times. For PVs with floating point values, the value may fluctuate around the `trip_point` for a while. If the value violates the `trip_point`, then momentarily goes back to an acceptable value, and back again to a violating value, it may not be desirable to send repeated, identical messages. To prevent this situation, the alarm callback will be called when the alarm condition is met **and** the callback was not called within the time specified by `alert_delay`.

6.2 Alarm Example

An epics Alarm is very easy to use. Here is an alarm set to print a message when a PV's value reaches a certain value:

```
from epics import Alarm, poll

def alertMe(pvname=None, char_value=None, **kw):
    print "Soup's on!  %s = %s" % (pvname, char_value)

my_alarm = Alarm(pvname = 'WaterTemperature.VAL',
                 comparison = '>',
                 callback = alertMe,
                 trip_point = 100.0,
                 alert_delay = 600)

while True:
    poll()
```

AUTO-SAVING: SIMPLE SAVE/RESTORE OF PVS

7.1 Overview

The `autosave` module provides simple save/restore functionality for PVs, similar to the `autosave` module in `synApps` for IOCs but (obviously) via Channel Access.

Request and Save file formats are designed to be compatible with `synApps` `autosave`.

Use of this module requires the `pyparsing` parser framework. The Debian/Ubuntu package is “python-pyparsing” The web site is <http://pyparsing.wikispaces.com/>

7.2 Examples

A simple example usign the `autosave` module:

```
import epics.autosave
# save values
epics.autosave.save_pvs("/tmp/my_request_file.req",
                        "/tmp/my_recent_save.sav")

# wait 30 seconds
time.sleep(30)

# restore those values back
epics.autosave.restore_pvs("/tmp/my_recent_save.sav")
```


WX: WXPYTHON WIDGETS FOR EPICS

The `wx` module of `epics` (that is, `epics.wx`) provides a set of wxPython classes for epics PVs. Most of these are derived from wxPython widgets, with special support added for epics PVs, especially regarding when to automatically update the widget based on a changing value for a PV.

Some examples of code that uses `pyepics` and wxPython are included in the `scripts` folder of the `pyepics` source distribution kit. In addition, there are some full-fledged applications using Epics and wxPython at [pyepics applications](#).

8.1 PV-aware Widgets

Several basic wxPython widgets have been extended so as to connect the widget with a corresponding *PV*. For example, setting the text value of a `PVTextCtrl` will forward that value to the epics *PV*, and if the *PV* is changed by externally, the value displayed in the widget will be automatically updated.

8.1.1 PVMixin

```
class wx.PVMixin ([pv=None[, pvname=None ]])
```

This is a mixin class for wx Controls with epics PVs: This connects to PV, and manages connection and callback events for the PV. It provides the following basic methods used by most of the PV<->widget classes below.

```
wx.SetPV (pv=None)
```

set the PV corresponding to the widget.

```
wx.Update (value=None)
```

set the widgets value from the PV's value. If value="None", the current value for the PV is used.

```
wx.GetValue (as_string=True)
```

return the PVs value.

```
wx.OnEpicsConnect ()
```

PV connection event handler.

```
wx.OnPVChange (value)
```

PV monitor (subscription) event handler. Must be overwritten for each widget type.

```
wx.GetEnumStrings ()
```

return enumeration strings for the PV

8.1.2 PVCtrlMixin

```
class wx.PVCtrlMixin (parent, pv=None, font=None, fg=None, bg=None, **kw)
```

This is a mixin class for wx Controls with epics PVs: This subclasses PVCtrlMixin and adds colour translations PV, and manages callback events for the PV.

Parameters

- **parent** – wx parent widget
- **pv** – epics.PV
- **font** – wx.Font for display
- **fg** – foreground colour
- **bg** – background colour

A class that inherits from this class **must** provide a method called `_SetValue`, which will set the contents of the corresponding widget when the PV's value changes.

In general, the widgets will automatically update when the PV changes. Where appropriate, setting the value with the widget will set the PV value.

8.1.3 PVText

```
class wx.PVText (parent, pv=None, font=None, fg=None, bg=None, minor_alarm="DARKRED",  
                 major_alarm="RED", invalid_alarm="ORANGERED", auto_units=False, units="",  
                 **kw)
```

derived from wx.StaticText and PVCtrlMixin, this is a StaticText widget whose value is set to the string representation of the value for the corresponding PV.

By default, the text colour will be overridden when the PV enters an alarm state. These colours can be modified (or disabled by being set to `None`) as part of the constructor.

“units” specifies a unit suffix (like ‘ A ’ or ‘ mm ’) to put after the text value whenever it is displayed.

Alternatively, “auto_units” means the control will automatically display the “EGU” units value from the PV, whenever it updates. If this value is set, “units” is ignored. A space is inserted between the value and the unit.

8.1.4 PVTextCtrl

```
class wx.PVTextCtrl (parent, pv=None, font=None, fg=None, bg=None, **kw)
```

derived from wx.TextCtrl and PVCtrlMixin, this is a TextCtrl widget whose value is set to the string representation of the value for the corresponding PV. Setting the value (hitting Return or Enter) of the widget will set the PV value.

8.1.5 PVFloatCtrl

```
class wx.PVFloatCtrl (parent, pv=None, font=None, fg=None, bg=None, **kw)
```

A special variation of a wx.TextCtrl that allows only floating point numbers, as associated with a double, float, or integer PV. Trying to type in a non-numerical value will be ignored. Furthermore, if a PV's limits can be determined, they will be used to limit the allowed range of input values. For a value that is within limits, the value will be *put* to the PV on return. Out-of-limit values will be highlighted in a different color.

8.1.6 PVBitmap

```
class wx.PVBitmap (parent, pv=None, bitmaps={}, defaultBitmap=None)
```

A Static Bitmap where the image is based on PV value.

If the bitmaps dictionary is set, it should be set as `PV.Value(Bitmap)` where particular bitmaps will be shown if the PV takes those certain values.

If you need to do any more complex or dynamic drawing, you may want to look at the OGL PV controls.

8.1.7 PVCheckBox

class `wx.PVCheckBox` (*self, parent, pv=None, on_value=1, off_value=0, **kw*)

Checkbox based on a binary PV value, both reads/writes the PV on changes. `on_value` and `off_value` are the specific values that are mapped to the checkbox.

There are multiple options for translating PV values to checkbox settings (from least to most complex):

- Use a PV with values 0 and 1
- Use a PV with values that convert via Python's own `bool(x)`
- Set `on_value` and `off_value` in the constructor
- Use `SetTranslations()` to set a dictionary for converting various PV values to booleans.

8.1.8 PVFloatSpin

class `wx.PVFloatSpin` (*parent, pv=None, deadTime=500, min_val=None, max_val=None, increment=1.0, digits=-1, **kw*)

A FloatSpin is a floating point spin control with buttons to increase and decrease the value by a particular increment. Arrow keys and page up/down can also be used (the latter changes the value by 10x the increment.)

PVFloatSpin is a special derivation that assigns a PV to the FloatSpin control. `deadTime` is the delay (in milliseconds) between when the user finishes typing a value and when the PV is set to it (to prevent half-typed numeric values being set.)

8.1.9 PVButton

PVButton (*parent, pv=None, pushValue=1, disablePV=None, disableValue=1, **kw*)

A `wx.Button` linked to a PV. When the button is pressed, 'pushValue' is written to the PV (useful for momentary PVs with `HIGH=` set.) Setting `disablePV` and `disableValue` will automatically cause the button to disable when that PV has a certain value.

8.1.10 PVRadioButton

class `wx.PVRadioButton` (*parent, pv=None, pvValue=None, **kw*)

A PVRadioButton is a radio button associated with a particular PV and one particular value.

Suggested for use in a group where all radio buttons are PVRadioButtons, and they all have a discrete value set.

8.1.11 PVComboBox

class `wx.PVComboBox` (*parent, pv=None, **kw*)

A ComboBox linked to a PV. Both reads/writes the combo value on changes.

8.1.12 PVEnumButtons

class wx.**PVEnumButtons** (*parent, pv=None, font=None, fg=None, bg=None, **kw*)

This will create a wx.Panel of buttons (a button bar), 1 for each enumeration state of an enum PV. The set of buttons will correspond to the current state of the PV

8.1.13 PVEnumChoice

class wx.**PVEnumChoice** (*parent, pv=None, font=None, fg=None, bg=None, **kw*)

This will create a dropdown list (a wx.Choice) with a list of enumeration states for an enum PV.

8.1.14 PVAlarm

class wx.**PVAlarm** (*parent, pv=None, font=None, fg=None, bg=None, trip_point=None, **kw*)

This will create a pop-up message (wx.MessageDialog) that is shown when the corresponding PV trips the alarm level.

8.1.15 PVCollapsiblePane

class wx.**PVCollapsiblePane** (*parent, pv=None, minor_alarm="DARKRED", major_alarm="RED",
invalid_alarm="ORANGERED", **kw*)

This is equivalent to wx.CollapsiblePane, except the label shown on the pane's "expansion button" comes from a PV.

The additional keyword arguments can be any of the other constructor arguments supported by wx.CollapsiblePane.

By default, the foreground colour of the pane button will be overridden when the PV enters an alarm state. On GTK, this means the colour of the triangular drop-down button but not the label text. These colours can be modified (or disabled by being set to None) as part of the constructor.

Supports the .SetTranslation() method, whose argument is a dictionary mapping PV values to display labels. If the PV value is not found in the dictionary, it will displayed verbatim as the label.

8.2 Decorators and other Utility Functions

wx.**DelayedEpicsCallback** ()

decorator to wrap an Epics callback in a wx.CallAfter, so that the wx and epics ca threads do not clash This also checks for dead wxPython objects (say, from a closed window), and remove callbacks to them.

wx.**EpicsFunction** ()

decorator to wrap function in a wx.CallAfter() so that Epics calls can be made in a separate thread, and asynchronously.

This decorator should be used for all code that mix calls to wx and epics

wx.**finalize_epics** ()

This function will finalize epics by calling `epics.ca.finalize_libca()`. It is recommended that this be added to any "close GUI" code, such as a method bound to `wx.EVT_CLOSE(self, self.onClose)`, where the function might look like this:



```
def onClose(self, event):
    finalize_epics()
    self.Destroy()
```

8.3 wxMotorPanel Widget

A dedicated wx Widget for Epics Motors is included in the `wx` module that provides an easy-to-use Motor panel that is similar to the normal MEDM window, but with a few niceties from the more sophisticated wx toolkit. This widget can be used simply as:

```
import wx
from epics.wx import MotorPanel
....
mymotor = MotorPanel(parent, 'XXX:m1')
```

A sample panel looks like this



Which shows from right to left: the motor description, an information message (blank most of the time), the readback value, the drive value, arrows to tweak the motor, and a drop-down combobox for tweak values, a “Stop” button and a “More” button. The panel has the following features:

- All controls are “live” and will respond to changes from other source.
- The values for the tweak values in the ComboBox are automatically generated from the precision and travel range of the motor.
- The entry box for the drive value will *only* accept numeric input, and will only set the drive value when hitting Enter or Return.
- The drive value will change to Red text on a Yellow background when the value in the box violates the motors (user) limits. If Enter or Return when the the displayed value violates the limit, the motor will not be moved, but the displayed value will be changed to the closest limit value.
- Pressing the “Stop” button will stop the motor (with the `.SPMG` field), and set the Info field to “Stopped”. The button label will change to “Go”, and the motor will not move until this button is pressed.

Finally, the “More” button will bring up a more complete form of Motor parameters that looks like:

Motor Details: 13XRM:m3 | theta |

13XRM:m3: (asynMotor)

Label

theta

units

degrees

Drive

	User	Dial	Raw	
High Limit	1500.0000	1500.0000		Stop
Readback	25.0000	25.0000	50000	Pause
Move	25.0000	25.0000	50000	Move
Low Limit	-1500.0000	-1500.0000		Go
Tweak	<div> <div><</div> <div>0.0100</div> <div>></div> </div>		<div>Enable</div> <div>Disable</div>	

Calibration

Mode:

Use

Set

Freeze Offset:

Variable

Direction:

Pos

Neg

Offset Value:

0.0000

Dynamics

	Normal	Backlash
Max Speed	20.0000	
Speed	20.0000	1.0000
Base Speed	1.0000	
Accel (s)	0.5000	0.2000
Backlash Distance		0.0000
Move Fraction		1.0000


Resolution

Motor Res	0.0005	Encoder Res	0.0005
Steps / Rev	2000	Units / Rev	1.0000
Precision	4		

Many such MotorPanels can be put in a vertical stack, as generated from the 'wx_motor.py' script in the scripts folder of the source distribution as:

```
~>python wx_motor.py XXX:m1 XXX:m2 XXX:m3 XXX:m4
```

will look like this:



8.4 OGL Classes

OGL is a graphics drawing library shipped with wxPython. Is it built around the concept of “shapes” which are added to “canvases” and can be moved, scrolled, zoomed, animated, etc.

There is a PVShapeMixin class which allows PV callback functionality to be added to any OGL Shape class, and there are also PVRectangle and PVCircle subclasses already created.

A recommended way to use these OGL classes is to make a static bitmap background for your display, place it in an OGL Canvas and then add an overlay of shapes which appear/disappear/resize/change colour based on the PV values.

8.4.1 PVShapeMixin

```
class wx.PVShapeMixin (self, pv=None, pvname=None)
```

Similar to PVMixin, this mixin should be added to any ogl.Shape subclass that needs PV callback support.

The main method is PVChanged(self, raw_value), which should be overridden in the subclass to provide specific processing based on the changed value.

There are also some built-in pieces of functionality. These are enough to do simple show/hide or change colour shape functionality, without needing to write specific code.

SetBrushTranslations(translations) allows setting a dict of PV Value -> wx.Brush mappings, which can be used to automatically repaint the shape foreground (fill) when the PV changes.

SetPenTranslations(translations) similar to brush translations, but the values are wx.Pen instances that are used to repaint the shape outline when the PV changes.

SetShownTranslations(translations) sets a dictionary of PV Value -> bool values which are used to show/hide the shape depending on the PV value, as it changes.

8.4.2 PVRectangle

```
class wx.PVRectangle (self, w, h, pv=None, pvname=None)
```

A PVCtrlMixin for the Rectangle shape class.

8.4.3 PVCircle

class wx.**PVCircle** (*self, diameter, pv=None, pvname=None*)
A PVCtrlMixin for the Circle shape class.

ADVANCED TOPIC WITH PYTHON

CHANNEL ACCESS

This chapter contains a variety of “usage notes” and implementation details that may help in getting the best performance from the pyepics module.

9.1 Working with waveform / array data

Though most EPICS Process Variables hold single values, PVs can hold array data from EPICS waveform records. These are always data of homogenous data type, and have a fixed maximum element count (defined when the waveform is created from the host EPICS process). These waveforms are most naturally mapped to Arrays from the [numpy module](#), and this is strongly encouraged.

9.1.1 Arrays without Numpy

If you have numpy installed, and use the default `as_numpy=True` in `ca.get()`, `pv.get()` or `epics.caget()`, you will get a numpy array for the value of a waveform PV. If you do *not* have numpy installed, or explicitly use `as_numpy=False` in a get request, you will get the raw C-like array reference from the Python [ctypes module](#). These objects are not normally meant for casual use, but are not too difficult to work with either. They can be easily converted to a simple Python list with something like:

```
>>> import epics
>>> epics.ca.HAS_NUMPY = False # turn numpy off for session
>>> p = epics.PV('XX:scan1.P1PA')
>>> p.get()
<lib.dbr.c_double_Array_500 object at 0x853980c>
>>> ldat = list(p.get())
```

Note that this conversion to a list can be very slow for large arrays.

9.1.2 Variable Length Arrays

While the maximum length of an array is fixed, the length of data you get back from a monitor, `ca.get()`, `pv.get()`, or `epics.caget()` may be shorter than the maximum length, reflecting the most recent data put to that PV. That is, if some process puts a smaller array to a PV than its maximum length, monitors on that PV may receive only the changed data. For example:

```
>>> import epics
>>> p = epics.PV('Py:double2k')
>>> print p
<PV 'Py:double2k', count=2048/2048, type=double, access=read/write>
>>> import numpy
>>> p.put(numpy.arange(10)/5.0)
>>> print p.get()
array([ 0. ,  0.2,  0.4,  0.6,  0.8,  1. ,  1.2,  1.4,  1.6,  1.8])
```

To be clear, the `pv.put()` above could be done in a separate process – the `pv.get()` is not using a value cached from the `pv.put()`.

This feature seems to depend on the record definition, and requires version 3.14.12.1 of Epics base or higher, and can be checked by comparing `ca.version()` with the string '4.13'.

9.1.3 Character Arrays

As noted in other sections, character waveforms can be used to hold strings longer than 40 characters, which a fundamental limit for native Epics strings. Character waveforms shorter than `ca.AUTOMONITOR_MAXLENGTH` can be turned into strings with an optional `as_string=True` to `ca.get()`, `pv.get()`, or `epics.caget()`. If you've defined a Epics waveform record as:

```
record(waveform, "$(P):filename") {
    field(DTYP, "Soft Channel")
    field(DESC, "file name")
    field(NELM, "128")
    field(FTVL, "CHAR")
}
```

Then you can use this record with:

```
>>> import epics
>>> pvname = 'PREFIX:filename.VAL'
>>> pv = epics.PV(pvname)
>>> print pv.info
....
>>> plain_val = pv.get()
>>> print plain_val
array([ 84,  58,  92, 120,  97, 115,  95, 117, 115, 101, 114,  92,  77,
        97, 114,  99, 104,  50,  48,  49,  48,  92,  70,  97, 115, 116,
        77,  97, 112])
>>> char_val = pv.get(as_string=True)
>>> print char_val
'T:\\xas_user\\March2010\\FastMap'
```

This example uses `pv.get()` but `ca.get()` is essentially equivalent, as its `as_string` parameter works exactly the same way.

9.2 Strategies for working with large arrays

EPICS Channels / Process Variables usually have values that can be stored with a small number of bytes. This means that their storage and transfer speeds over real networks is not a significant concern. However, some Process Variables can store much larger amounts of data (say, several megabytes) which means that some of the assumptions about dealing with Channels / PVs may need reconsideration.

When using PVs with large array sizes (here, I'll assert that *large* means more than a few thousand elements), it is necessary to make sure that the environmental variable `EPICS_CA_MAX_ARRAY_BYTES` is suitably set. Unfortunately, this represents a pretty crude approach to memory management within Epics for handling array data as it is used not only sets how large an array the client can accept, but how much memory will be allocated on the server. In addition, this value must be set prior to using the CA library – it cannot be altered during the running of a CA program.

Normally, the default value for `EPICS_CA_MAX_ARRAY_BYTES` is 16384 (16k, and it turns out that you cannot set it smaller than this value!). As Python is used for clients, generally running on workstations or servers with sufficient memory, this default value is changed to 2^{24} , or 16Mb) when `epics.ca` is initialized. If the environmental variable `EPICS_CA_MAX_ARRAY_BYTES` has not already been set.

The other main issue for PVs holding large arrays is whether they should be automatically monitored. For PVs holding scalar data or small arrays, any penalty for automatically monitoring these variables (that is, causing network traffic every time a PV changes) is a small price to pay for being assured that the latest value is always available. As arrays get larger (as for data streams from Area Detectors), it is less obvious that automatic monitoring is desirable.

The Python `epics.ca` module defines a variable `ca.AUTOMONITOR_MAXLENGTH` which controls whether array PVs are automatically monitored. The default value for this variable is 65536, but can be changed at runtime. Arrays with fewer elements than `ca.AUTOMONITOR_MAXLENGTH` will be automatically monitored, unless explicitly set, and arrays larger than `AUTOMONITOR_MAXLENGTH` will not be automatically monitored unless explicitly set. Auto-monitoring of PVs can be explicitly set with

```
>>> pv2 = epics.PV('ScalerPV', auto_monitor=True)
>>> pv1 = epics.PV('LargeArrayPV', auto_monitor=False)
```

9.2.1 Example handling Large Arrays

Here is an example reading data from an `EPICS areaDetector`, as if it were an image from a digital camera. This uses the `Python Imaging Library` for much of the image processing:

```
>>> import epics
>>> import Image
>>> pvname = '13IDCPS1:image1:ArrayData'
>>> img_pv = epics.PV(pvname)
>>>
>>> raw_image = img_pv.get()
>>> im_mode = 'RGB'
>>> im_size = (1360, 1024)
>>> img = Image.frombuffer(im_mode, im_size, raw_image,
>>>                        'raw', im_mode, 0, 1)
>>> img.show()
```

The result looks like this (taken with a Prosilica GigE camera):



A more complete application for reading and displaying image from Epics Area Detectors is included at <http://github.com/pyepics/epicsapps/>.

9.3 Using Python Threads

An important feature of the PyEpics package is that it can be used with Python threads, as Epics 3.14 supports threads for client code. Even in the best of cases, working with threads can be somewhat tricky and lead to unexpected behavior, and the Channel Access library adds a small level of complication for using CA with Python threads. The result is that some precautions may be in order when using PyEpics and threads. This section discusses the strategies for using threads with PyEpics.

First, to use threads with Channel Access, you must have `epics.ca.PREEMPTIVE_CALLBACK = True`. This is the default value, but if `epics.ca.PREEMPTIVE_CALLBACK` has been set to `False`, threading will not work.

Second, if you are using PV objects and not making heavy use of the `ca` module (that is, not making and passing around chids), then the complications below are mostly hidden from you. If you're writing threaded code, it's probably a good idea to read this just to understand what the issues are.

9.3.1 Channel Access Contexts

The Channel Access library uses a concept of *contexts* for its own thread model, with contexts holding sets of threads as well as Channels and Process Variables. For non-threaded work, a process will use a single context that is initialized prior doing any real CA work (done in `ca.initialize_libca()`). In a threaded application, each new thread

begins with a new, uninitialized context that must be initialized or replaced. Thus each new python thread that will interact with CA must either explicitly create its own context with `ca.create_context()` (and then, being a good citizen, destroy this context as the thread ends with `ca.destroy_context()`) or attach to an existing context.

The generally recommended approach is to use a single CA context throughout an entire process and have each thread attach to the first context created (probably from the main thread). This avoids many potential pitfalls (and crashes), and can be done fairly simply. It is the default mode when using PV objects.

The most explicit use of contexts is to put `epics.ca.create_context()` at the start of each function call as a thread target, and `epics.ca.destroy_context()` at the end of each thread. This will cause all the activity in that thread to be done in its own context. This works, but means more care is needed, and so is not the recommended.

The best way to attach to the initially created context is to call `epics.ca.use_initial_context()` before any other CA calls in each function that will be called by `Thread.run()`. Equivalently, you can add a `withInitialContext()` decorator to the function. Creating a PV object will implicitly do this for you, as long as it is your first CA action in the function. Each time you do a `PV.get()` or `PV.put()` (or a few other methods), it will also check that the initial context is being used.

Of course, this approach requires CA to be initialized already. Doing that *in the main thread* is highly recommended. If it happens in a child thread, that thread must exist for all CA work, so either the life of the process or with great care for processes that do only some CA calls. If you are writing a threaded application in which the first real CA calls are inside a child thread, it is recommended that you initialize CA in the main thread,

As a convenience, the `CAThread` in the `ca` module is a very thin wrapper around the standard `threading.Thread` which adding a call of `epics.ca.use_initial_context()` just before your threaded function is run. This allows your target functions to not explicitly set the context, but still ensures that the initial context is used in all functions.

9.3.2 How to work with CA and Threads

Summarizing the discussion above, to use threads you must use run in `PREEMPTIVE_CALLBACK` mode. Furthermore, it is recommended that you use a single context, and that you initialize CA in the main program thread so that your single CA context belongs to the main thread. Using PV objects exclusively makes this easy, but it can also be accomplished relatively easily using the lower-level ca interface. The options for using threads (in approximate order of reliability) are then:

1. use PV objects for threading work.
2. use `CAThread` instead of `Thread` for threads that will use CA calls.
3. put `epics.ca.use_initial_context()` at the top of all functions that might be a `Thread` target function, or decorate them with `withInitialContext()` decorator, `@withInitialContext`.
4. use `epics.ca.create_context()` at the top of all functions that are inside a new thread, and be sure to put `epics.ca.destroy_context()` at the end of the function.
5. ignore this advise and hope for the best. If you're not creating new PVs and only reading values of PVs created in the main thread inside a child thread, you may not see a problems, at least not until you try to do something fancier.

9.3.3 Thread Examples

This is a simplified version of test code using Python threads. It is based on code originally from Friedrich Schotte, NIH, and included as `thread_test.py` in the `tests` directory of the source distribution.

In this example, we define a `run_test` procedure which will create PVs from a supplied list, and monitor these PVs, printing out the values when they change. Two threads are created and run concurrently, with overlapping PV lists, though one thread is run for a shorter time than the other.

```
"""This script tests using EPICS CA and Python threads together
Based on code from Friedrich Schotte, NIH, modified by Matt Newville
19-Apr-2010
"""
import time
from sys import stdout
from threading import Thread
import epics

from pvnames import updating_pvlist
pvlist_a = updating_pvlist[:-1]
pvlist_b = updating_pvlist[1:]

def run_test(runtime=1, pvnames=None, run_name='thread c'):
    msg = '-> thread "%s" will run for %.3f sec, monitoring %s\n'
    stdout.write(msg % (run_name, runtime, pvnames))
    def onChanges(pvname=None, value=None, char_value=None, **kw):
        stdout.write('    %s = %s (%s)\n' % (pvname, char_value, run_name))
        stdout.flush()

    # epics.ca.use_initial_context() # epics.ca.create_context()
    start_time = time.time()
    pvs = [epics.PV(pvn, callback=onChanges) for pvn in pvnames]

    while time.time()-start_time < runtime:
        time.sleep(0.1)

    [p.clear_callbacks() for p in pvs]
    stdout.write('Completed Thread %s\n' % (run_name))

stdout.write("First, create a PV in the main thread:\n")
p = epics.PV(updating_pvlist[0])

stdout.write("Run 2 Background Threads simultaneously:\n")
th1 = Thread(target=run_test, args=(3, pvlist_a, 'A'))
th1.start()

th2 = Thread(target=run_test, args=(6, pvlist_b, 'B'))
th2.start()

th2.join()
th1.join()
stdout.write('Done\n')
```

In light of the long discussion above, a few remarks are in order: This code uses the standard Thread library and explicitly calls `epics.ca.use_initial_context()` prior to any CA calls in the target function. Also note that the `run_test()` function is first called from the main thread, so that the initial CA context does belong to the main thread. Finally, the `epics.ca.use_initial_context()` call in `run_test()` above could be replaced with `epics.ca.create_context()`, and run OK.

The output from this will look like:

```
First, create a PV in the main thread:
Run 2 Background Threads simultaneously:
-> thread "A" will run for 3.000 sec, monitoring ['Py:ao1', 'Py:ai1', 'Py:long1']
-> thread "B" will run for 6.000 sec, monitoring ['Py:ai1', 'Py:long1', 'Py:ao2']
Py:ao1 = 8.3948 (A)
Py:ai1 = 3.14 (B)
Py:ai1 = 3.14 (A)
```

```

Py:ao1 = 0.7404 (A)
Py:ai1 = 4.07 (B)
Py:ai1 = 4.07 (A)
Py:long1 = 3 (B)
Py:long1 = 3 (A)
Py:ao1 = 13.0861 (A)
Py:ai1 = 8.49 (B)
Py:ai1 = 8.49 (A)
Py:ao2 = 30 (B)
Completed Thread A
Py:ai1 = 9.42 (B)
Py:ao2 = 30 (B)
Py:long1 = 4 (B)
Py:ai1 = 3.35 (B)
Py:ao2 = 31 (B)
Py:ai1 = 4.27 (B)
Py:ao2 = 31 (B)
Py:long1 = 5 (B)
Py:ai1 = 8.20 (B)
Py:ao2 = 31 (B)
Completed Thread B
Done

```

Note that while both threads *A* and *B* are running, a callback for the PV *Py:ai1* is generated in each thread.

Note also that the callbacks for the PVs created in each thread are **explicitly cleared** with:

```
[p.clear_callbacks() for p in pvs]
```

Without this, the callbacks for thread *A* will persist even after the thread has completed!

9.4 The wait and timeout options for `get()`, `ca.get_complete()`

The `get` functions, `epics.caget()`, `pv.get()` and `ca.get()` all ask for data to be transferred over the network. For large data arrays or slow networks, this can take a noticeable amount of time. For PVs that have been disconnected, the `get` call will fail to return a value at all. For this reason, these functions all take a `timeout` keyword option. The lowest level `ca.get()` also has a `wait` option, and a companion function `ca.get_complete()`. This section describes the details of these.

If you're using `epics.caget()` or `pv.get()` you can supply a timeout value. If the value returned is `None`, then either the PV has truly disconnected or the timeout passed before receiving the value. If the `get` is incomplete in this way, a subsequent `epics.caget()` or `pv.get()` may actually complete and receive the value.

At the lowest level (which `pv.get()` and `epics.caget()` use), `ca.get()` issues a get-request with an internal callback function (that is, it calls `libca.ca_array_get_callback()` with a pre-defined callback function). With `wait=True` (the default), `ca.get()` then waits up to the timeout or until the CA library calls the specified callback function. If the callback has been called, the value can then be converted and returned.

If the callback is not called in time or if `wait=False` is used but the PV is connected, the callback will be called eventually, and simply waiting (or using `ca.pend_event()` if `ca.PREEMPTIVE_CALLBACK` is `False`) may be sufficient for the data to arrive. Under this condition, you can call `ca.get_complete()`, which will NOT issue a new request for data to be sent, but wait (for up to a timeout time) for the previous get request to complete.

`ca.get_complete()` will return `None` if the timeout is exceeded or if there is not an "incomplete get" that it can wait to complete. Thus, you should use the return value from `ca.get_complete()` with care.

Note that `pv.get()` (and so `epics.caget()`) will normally rely on the PV value to be filled in automatically by monitor callbacks. If monitor callbacks are disabled (as is done for large arrays and can be turned off) or if the monitor

hasn't been called yet, `pv.get()` will check whether it should call `ca.get()` or `ca.get_complete()`.

If not specified, the timeout for `ca.get_complete()` (and all other get functions) will be set to:

```
timeout = 0.5 + log10(count)
```

Again, that's the maximum time that will be waited, and if the data is received faster than that, the `get` will return as soon as it can.

9.5 Strategies for connecting to a large number of PVs

Occasionally, you may find that you need to quickly connect to a large number of PVs, say to write values to disk. The most straightforward way to do this, say:

```
import epics

pvnamelist = read_list_pvs()
pv_vals = {}
for name in pvnamelist:
    pv = epics.PV(name)
    pv_vals[name] = pv.get()
```

does incur some small performance penalty. As shown below, the penalty is generally pretty small in absolute terms, but can be noticeable when you are connecting to a large number (say, more than 100) PVs at once.

The cause for the penalty, and its remedy, are two-fold. First, a *PV* object automatically use connection and event callbacks. Normally, these are advantages, as you don't need to explicitly deal with them. But, internally, they do pause for network responses using `ca.pend_event()` and these pauses can add up. Second, the `ca.get()` also pauses for network response, so that the returned value actually contains the latest data right away, as discussed in the previous section.

The remedies are to

1. not use connection or event callbacks.
2. not explicitly wait for values to be returned for each `get()`.

A more complicated but faster approach relies on a carefully-tuned use of the CA library, and would be the following:

```
from epics import ca

pvnamelist = read_list_pvs()

pvdata = {}
for name in pvnamelist:
    chid = ca.create_channel(name, connect=False, auto_cb=False) # note 1
    pvdata[name] = (chid, None)

for name, data in pvdata.items():
    ca.connect_channel(data[0])
ca.poll()
for name, data in pvdata.items():
    ca.get(data[0], wait=False) # note 2

ca.poll()
for name, data in pvdata.items():
    val = ca.get_complete(data[0])
    pvdata[name][1] = val
```

```
for name, data in pvdata.items():
    print name, data[1]
```

The code here probably needs detailed explanation. The first thing to notice is that this is using the *ca* level, not *PV* objects. Second (Note 1), the *connect=False* and *auto_cb=False* options to `ca.create_channel()`. These respectively tell `ca.create_channel()` to not wait for a connection before returning, and to not automatically assign a connection callback. Normally, these are not what you want, as you want a connected channel and to know if the connection state changes. But we're aiming for maximum speed here, so we avoid these.

We then explicitly call `ca.connect_channel()` for all the channels. Next (Note 2), we tell the CA library to request the data for the channel without waiting around to receive it. The main point of not having `ca.get()` wait for the data for each channel as we go is that each data transfer takes time. Instead we request data to be sent in a separate thread for all channels without waiting. Then we do wait by calling `ca.poll()` once and only once, (not `len(channels)` times!). Finally, we use the `ca.get_complete()` method to convert the data that has now been received by the companion thread to a python value.

How much faster is the more explicit method? In my tests, I used 20,000 PVs, all scalar values, all actually connected, and all on the same subnet as the test client, though on a mixture of several vxWorks and linux IOCs. I found that the simplest, obvious approach as above took around 12 seconds to read all 20,000 PVs. Using the *ca* layer with connection callbacks and a normal call to `ca.get()` also took about 12 seconds. The method without connection callbacks and with delayed unpacking above took about 2 seconds to read all 20,000 PVs.

Is that performance boost from 12 to 2 seconds significant? If you're writing a script that is intended to run once, fetch a large number of PVs and get their values (say, an auto-save script that runs on demand), then the boost is definitely significant. On the other hand, if you're writing a long running process or a process that will retain the PV connections and get their values multiple times, the difference in start-up speed is less significant. For a long running auto-save script that periodically writes out all the PV values, the "obvious" way using automatically monitored PVs may be much *better*, as the time for the initial connection is small, and the use of event callbacks will reduce network traffic for PVs that don't change between writes.

Note that the tests also show that, with the simplest approach, 1,000 PVs should connect and receive values in under 1 second. Any application that is sure it needs to connect to PVs faster than that rate will want to do careful timing tests. Finally, note also that the issues are not really a classic *python is slow compared to C* issue, but rather a matter of how much pausing with `ca.poll()` one does to make sure values are immediately useful.

9.6 time.sleep() or epics.poll()?

In order for a program to communicate with Epics devices, it needs to allow some time for this communication to happen. With `ca.PREEMPTIVE_CALLBACK` set to `True`, this communication will be handled in a thread separate from the main Python thread. This means that CA events can happen at any time, and `ca.pend_event()` does not need to be called to explicitly allow for event processing.

Still, some time must be released from the main Python thread on occasion in order for events to be processed. The simplest way to do this is with `time.sleep()`, so that an event loop can simply be:

```
>>> while True:
>>>     time.sleep(0.001)
```

Unfortunately, the `time.sleep()` method is not a very high-resolution clock, with typical resolutions of 1 to 10 ms, depending on the system. Thus, even though events will be asynchronously generated and epics with pre-emptive callbacks does not *require* `ca.pend_event()` or `ca.poll()` to be run, better performance may be achieved with an event loop of:

```
>>> while True:
>>>     epics.poll(evt=1.e-5, iot=0.1)
```

as the loop will be run more often than using `time.sleep()`.

PYTHON MODULE INDEX

a

alarm, [53](#)
autosave, [55](#)

c

ca, [29](#)

d

device, [45](#)

e

epics, [9](#)

m

motor, [46](#)

p

pv, [15](#)

w

wx, [57](#)

INDEX

Symbols

`_cache` (in module `ca`), 33
`_pvs` (in module `device`), 46
`_unpack` (in module `ca`), 39

A

`access` (in module `pv`), 19
`access` (in module `ca`), 33
`add_callback` (in module `device`), 45
`add_callback` (in module `pv`), 17
`add_pv` (in module `device`), 45
`Alarm` (class in `alarm`), 53
`alarm` (module), 53
`attach_context` (in module `ca`), 31
`AUTOMONITOR_MAXLENGTH` (in module `ca`), 30
`autosave` (module), 55

C

`ca` (module), 29
`ca_add_exception_event` (in module `ca`), 40
`ca_add_fd_registration` (in module `ca`), 40
`ca_client_status` (in module `ca`), 40
`ca_dump_dbr` (in module `ca`), 41
`ca_puser` (in module `ca`), 40
`ca_replace_access_rights_event` (in module `ca`), 40
`ca_set_puser` (in module `ca`), 40
`ca_SEVCHK` (in module `ca`), 40
`ca_signal` (in module `ca`), 41
`ca_test_event` (in module `ca`), 41
`caget` (in module `epics`), 9
`cainfo` (in module `epics`), 11
`callbacks` (in module `pv`), 19
`camonitor` (in module `epics`), 11
`camonitor_clear` (in module `epics`), 12
`caput` (in module `epics`), 11
`CAThread` (class in `ca`), 41
`ChannelAccessException`, 38
`char_value` (in module `pv`), 18
`check_limits` (in module `motor`), 48
`clear_callback` (in module `motor`), 50
`clear_callbacks` (in module `pv`), 18

`clear_channel` (in module `ca`), 32
`clear_subscription` (in module `ca`), 35
`client_status` (in module `ca`), 31
`connect` (in module `pv`), 17
`connect_channel` (in module `ca`), 32
`connection_callbacks` (in module `pv`), 20
`context_create` (in module `ca`), 31
`context_destroy` (in module `ca`), 31
`count` (in module `pv`), 19
`create_channel` (in module `ca`), 31
`create_context` (in module `ca`), 31
`create_subscription` (in module `ca`), 35
`current_context` (in module `ca`), 31

D

`DEFAULT_CONNECTION_TIMEOUT` (in module `ca`), 30
`DEFAULT_SUBSCRIPTION_MASK` (in module `ca`), 35
`DelayedEpicsCallback` (in module `wx`), 60
`destroy_context` (in module `ca`), 31
`detach_context` (in module `ca`), 31
`Device` (class in `device`), 45
`device` (module), 45
`disconnect` (in module `pv`), 17

E

`element_count` (in module `ca`), 32
`enum_strs` (in module `pv`), 19
`epics` (module), 9
`EpicsFunction` (in module `wx`), 60

F

`field_type` (in module `ca`), 32
`finalize_epics` (in module `wx`), 60
`flush_io` (in module `ca`), 31
`ftype` (in module `pv`), 18

G

`get` (in module `ca`), 33
`get` (in module `device`), 45
`get` (in module `motor`), 48

get() (in module pv), 16
get_complete() (in module ca), 34
get_ctrlvars() (in module ca), 36
get_ctrlvars() (in module pv), 17
get_enum_strings() (in module ca), 36
get_position() (in module motor), 49
get_precision() (in module ca), 36
get_pv() (in module motor), 50
get_severity() (in module ca), 35
get_timestamp() (in module ca), 35
get_timevars() (in module ca), 36
GetEnumStrings() (in module wx), 57
GetValue() (in module wx), 57

H

host (in module pv), 19
host_name() (in module ca), 32

I

info (in module pv), 19
initialize_libca() (in module ca), 30
isConnected() (in module ca), 33

L

lower_alarm_limit (in module pv), 19
lower_ctrl_limit (in module pv), 19
lower_disp_limit (in module pv), 19
lower_warning_limit (in module pv), 19

M

message() (in module ca), 31
Motor (class in motor), 47
motor (module), 46
move() (in module motor), 49

N

name() (in module ca), 32
nelm (in module pv), 19

O

OnEpicsConnect() (in module wx), 57
OnPVChange() (in module wx), 57

P

pend_event() (in module ca), 31
pend_io() (in module ca), 31
poll() (in module ca), 31
poll() (in module pv), 17
precision (in module pv), 19
PREEMPTIVE_CALLBACK (in module ca), 30
promote_type() (in module ca), 33
put() (in module ca), 34
put() (in module device), 45

put() (in module motor), 48
put() (in module pv), 16
put_complete (in module pv), 19
PV (class in pv), 15
pv (module), 15
PV() (in module device), 45
PVALarm (class in wx), 60
PVBitmap (class in wx), 58
PVCheckBox (class in wx), 59
PVCircle (class in wx), 64
PVCollapsiblePane (class in wx), 60
PVComboBox (class in wx), 59
PVCtrlMixin (class in wx), 57
PVEnumButtons (class in wx), 60
PVEnumChoice (class in wx), 60
PVFloatCtrl (class in wx), 58
PVFloatSpin (class in wx), 59
PVMixin (class in wx), 57
PVRadioButton (class in wx), 59
PVRectangle (class in wx), 63
PVShapeMixin (class in wx), 63
PVText (class in wx), 58
PVTextCtrl (class in wx), 58
PySEVCHK() (in module ca), 38

R

read_access (in module pv), 19
read_access() (in module ca), 32
read_state() (in module device), 46
remove_callback() (in module pv), 18
replace_printf_handler() (in module ca), 31
restore_state() (in module device), 46
run_callback() (in module pv), 18
run_callbacks() (in module pv), 18

S

save_state() (in module device), 46
set_callback() (in module motor), 50
set_position() (in module motor), 49
SetPV() (in module wx), 57
severity (in module pv), 19
sg_block() (in module ca), 36
sg_create() (in module ca), 36
sg_delete() (in module ca), 36
sg_get() (in module ca), 36
sg_put() (in module ca), 37
sg_reset() (in module ca), 37
sg_test() (in module ca), 37
show_cache() (in module ca), 33
show_info() (in module motor), 50
state() (in module ca), 32
status (in module pv), 18

T

timestamp (in module pv), [19](#)
tweak() (in module motor), [49](#)
type (in module pv), [18](#)

U

units (in module pv), [19](#)
Update() (in module wx), [57](#)
upper_alarm_limit (in module pv), [19](#)
upper_ctrl_limit (in module pv), [19](#)
upper_disp_limit (in module pv), [19](#)
upper_warning_limit (in module pv), [19](#)
use_initial_context() (in module ca), [31](#)

V

value (in module pv), [18](#)
version() (in module ca), [31](#)

W

wait_for_connection() (in module pv), [17](#)
withCA() (in module ca), [39](#)
withCHID() (in module ca), [39](#)
withConnectedCHID() (in module ca), [39](#)
within_limits() (in module motor), [48](#)
withInitialContext() (in module ca), [39](#)
withSEVCHK() (in module ca), [38](#)
write_access (in module pv), [19](#)
write_access() (in module ca), [32](#)
write_state() (in module device), [46](#)
wx (module), [57](#)